

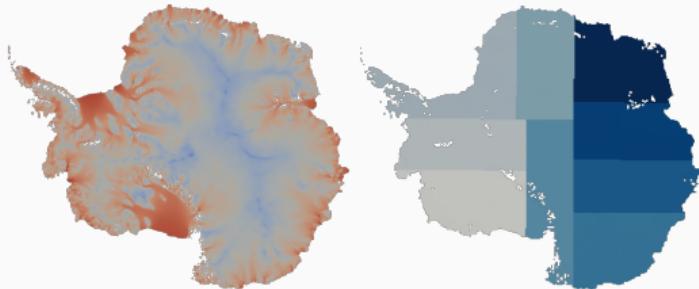
Domain decomposition for neural networks

Alexander Heinlein¹

IGHASC: Indo-German Workshop on Hardware-aware Scientific Computing, Heidelberg, Germany,
October 28-30, 2024

¹Delft University of Technology

Domain Decomposition Methods



Images based on Heinlein, Perego, Rajamanickam (2022)

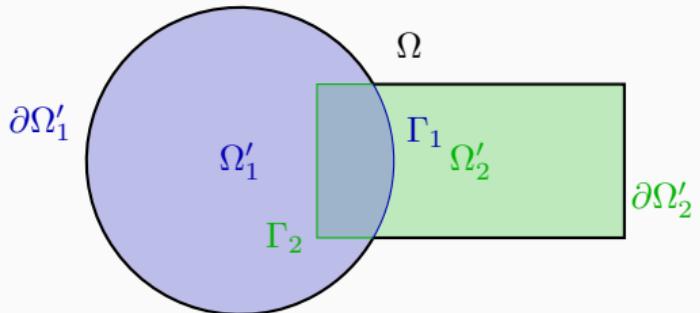
Historical remarks: The **alternating Schwarz method** is the earliest **domain decomposition method (DDM)**, which has been invented by **H. A. Schwarz** and published in **1870**:

- Schwarz used the algorithm to establish the **existence of harmonic functions** with prescribed boundary values on **regions with non-smooth boundaries**.

Idea

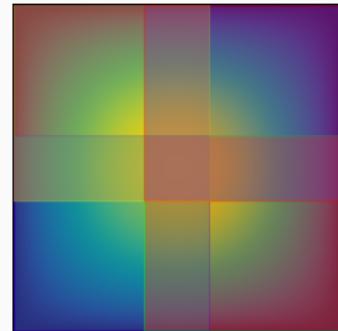
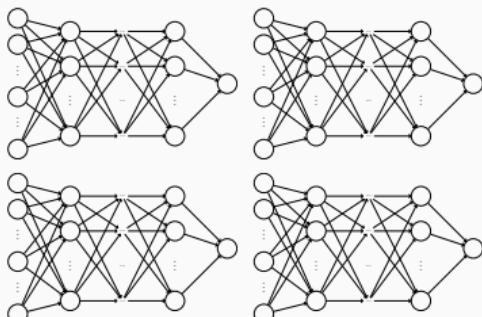
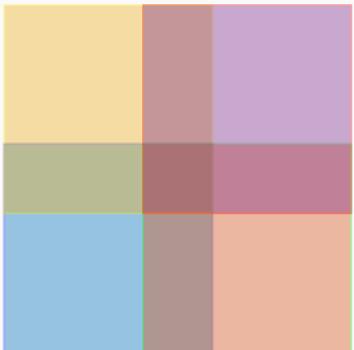
Decomposing a large **global problem** into smaller **local problems**:

- Better robustness** and **scalability** of numerical solvers
- Improved computational efficiency**
- Introduce **parallelism**



Domain Decomposition for Neural Networks

I)



II)

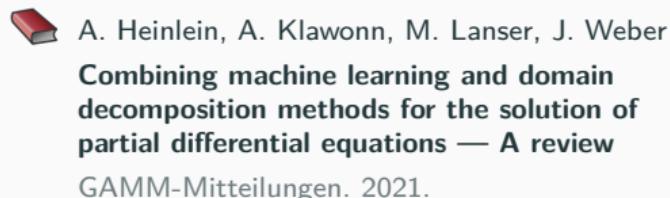


Domain Decomposition Methods and Machine Learning – Literature

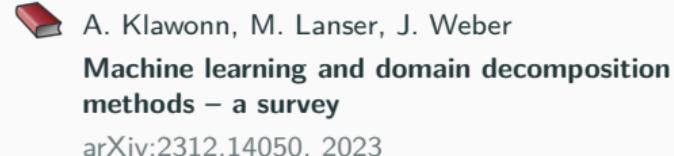
A non-exhaustive literature overview:

- Machine Learning for adaptive BDDC, FETI–DP, and AGDSW: Heinlein, Klawonn, Lanser, Weber (2019, 2020, 2021, 2021, 2021, 2022); Klawonn, Lanser, Weber (2024)
- cPINNs, XPINNs: Jagtap, Kharazmi, Karniadakis (2020); Jagtap, Karniadakis (2020)
- Classical Schwarz iteration for PINNs or DeepRitz (D3M, DeepDDM, etc):: Li, Tang, Wu, and Liao (2019); Li, Xiang, Xu (2020); Mercier, Gratton, Boudier (arXiv 2021); Dolean, Heinlein, Mercier, Gratton (subm. 2024 / arXiv:2408.12198); Li, Wang, Cui, Xiang, Xu (2023); Sun, Xu, Yi (arXiv 2022, arXiv 2023); Kim, Yang (2022, arXiv 2023)
- FBPINNs, FBKANs: Moseley, Markham, and Nissen-Meyer (2023); Dolean, Heinlein, Mishra, Moseley (2024, 2024); Heinlein, Howard, Beecroft, Stinis (acc. 2024 / arXiv:2401.07888); Howard, Jacob, Murphy, Heinlein, Stinis (arXiv:2406.19662)
- DDMs for CNNs: Gu, Zhang, Liu, Cai (2022); Lee, Park, Lee (2022); Klawonn, Lanser, Weber (2024); Verburg, Heinlein, Cyr (subm. 2024)

An overview of the state-of-the-art in early 2021:



An overview of the state-of-the-art in late 2023:



Outline

1 Multilevel domain decomposition-based architectures for physics-informed neural networks

Based on joint work with

Victorita Dolean

Ben Moseley and Siddhartha Mishra

(Eindhoven University of Technology)
(ETH Zürich)

2 Domain Decomposition for Convolutional Neural Networks

Based on joint work with

Eric Cyr

Corné Verburg

(Sandia National Laboratories)
(Delft University of Technology)

Multilevel domain decomposition-based architectures for physics-informed neural networks

Physics-Informed Neural Networks (PINNs) – Idea

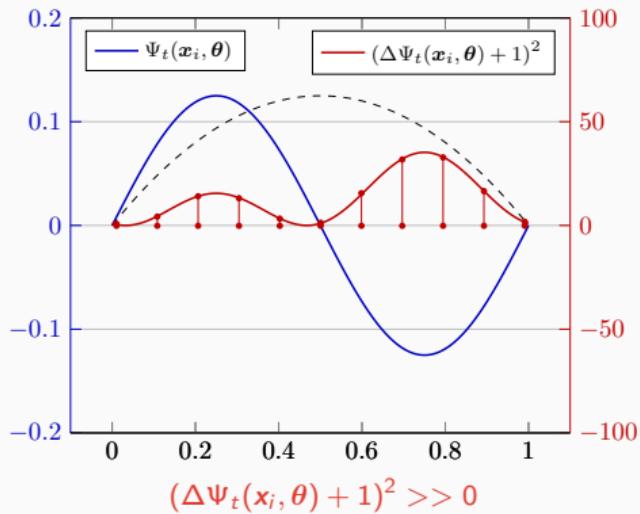
In [Lagaris et al. \(1998\)](#), the authors solve the boundary value problem

$$-\Delta \Psi_t(x, \theta) = 1 \text{ on } [0, 1],$$

$$\Psi_t(0, \theta) = \Psi_t(1, \theta) = 0,$$

via a **collocation approach**:

$$\min_{\theta} \sum_{x_i} (\Delta \Psi_t(x_i, \theta) + 1)^2$$

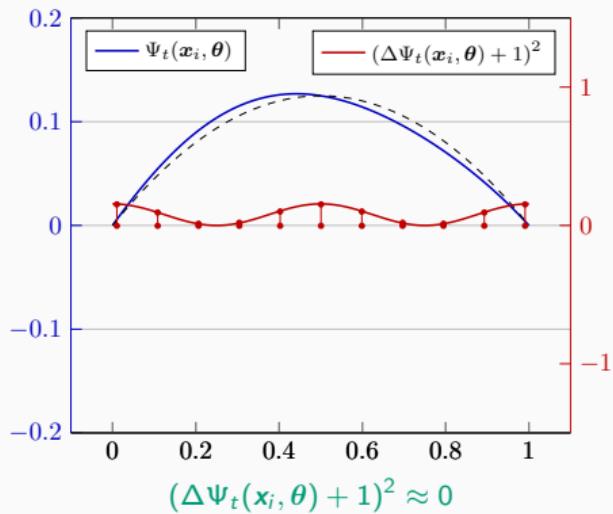


Boundary conditions ...

... can be enforced explicitly via the ansatz:

$$\Psi_t(x, \theta) = A(x) + F(x, \text{NN}(x, \theta))$$

- A satisfies the boundary conditions
- F does not contribute to the boundary conditions



Physics-Informed Neural Networks (PINNs)

In the physics-informed neural network (PINN) approach introduced by [Raissi et al. \(2019\)](#), a neural network is employed to discretize a partial differential equation

$$n[u] = f, \quad \text{in } \Omega.$$

PINNs use a **hybrid loss function**:

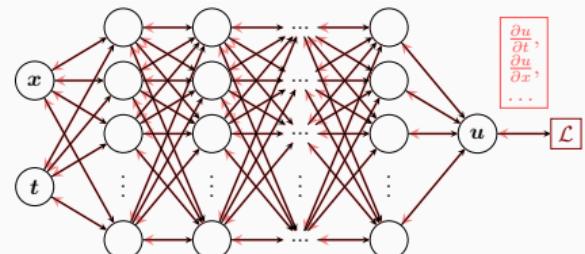
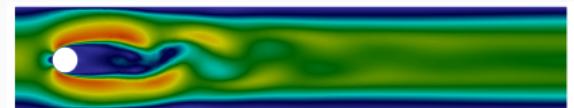
$$\mathcal{L}(\theta) = \omega_{\text{data}} \mathcal{L}_{\text{data}}(\theta) + \omega_{\text{PDE}} \mathcal{L}_{\text{PDE}}(\theta),$$

where ω_{data} and ω_{PDE} are **weights** and

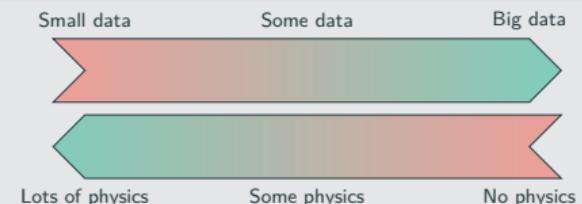
$$\mathcal{L}_{\text{data}}(\theta) = \frac{1}{N_{\text{data}}} \sum_{i=1}^{N_{\text{data}}} (u(\hat{x}_i, \theta) - u_i)^2,$$

$$\mathcal{L}_{\text{PDE}}(\theta) = \frac{1}{N_{\text{PDE}}} \sum_{i=1}^{N_{\text{PDE}}} (n[u](x_i, \theta) - f(x_i))^2.$$

See also [Dissanayake and Phan-Thien \(1994\)](#); [Lagaris et al. \(1998\)](#).



Hybrid loss



Advantages

- "Meshfree"
- Small data
- Generalization properties
- High-dimensional problems
- Inverse and parameterized problems

Drawbacks

- Training cost and robustness
- Convergence not well-understood
- Difficulties with scalability and multi-scale problems

- Known solution values can be included in $\mathcal{L}_{\text{data}}$
- Initial and boundary conditions are also included in $\mathcal{L}_{\text{data}}$

Theoretical Result for PINNs

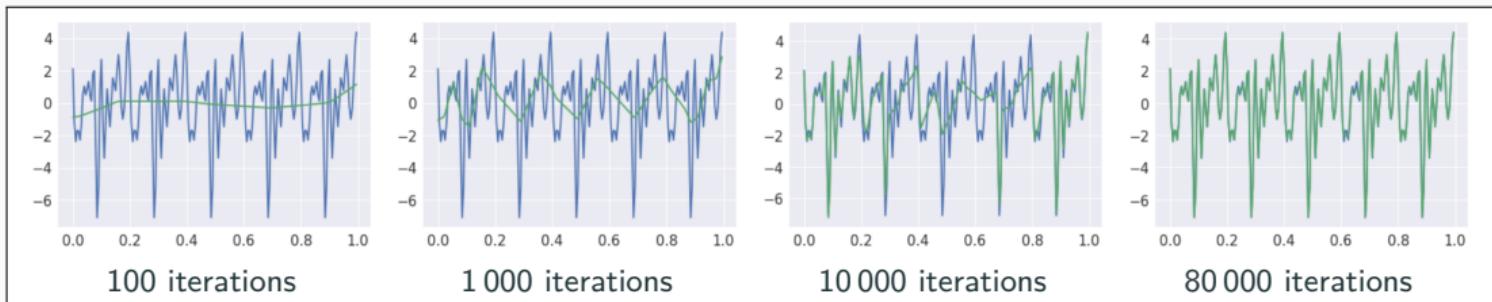
Estimate of the generalization error (Mishra and Molinaro (2022))

The generalization error (or total error) satisfies

$$\mathcal{E}_G \leq C_{\text{PDE}} \mathcal{E}_{\mathcal{T}} + C_{\text{PDE}} C_{\text{quad}}^{1/p} N^{-\alpha/p}$$

- $\mathcal{E}_G = \mathcal{E}_G(\mathbf{X}, \theta) := \|\mathbf{u} - \mathbf{u}^*\|_V$ **general. error** (V Sobolev space, \mathbf{X} training data set)
- $\mathcal{E}_{\mathcal{T}}$ **training error** (l^p loss of the residual of the PDE)
- N **number of the training points** and α **convergence rate of the quadrature**
- C_{PDE} and C_{quad} **constants** depending on the **PDE, quadrature, and neural network**

Rule of thumb: “As long as the PINN is **trained well**, it also **generalizes well**”



Rahaman et al., *On the spectral bias of neural networks*, ICML (2019)

Motivation – Some Observations on the Performance of PINNs

Solve

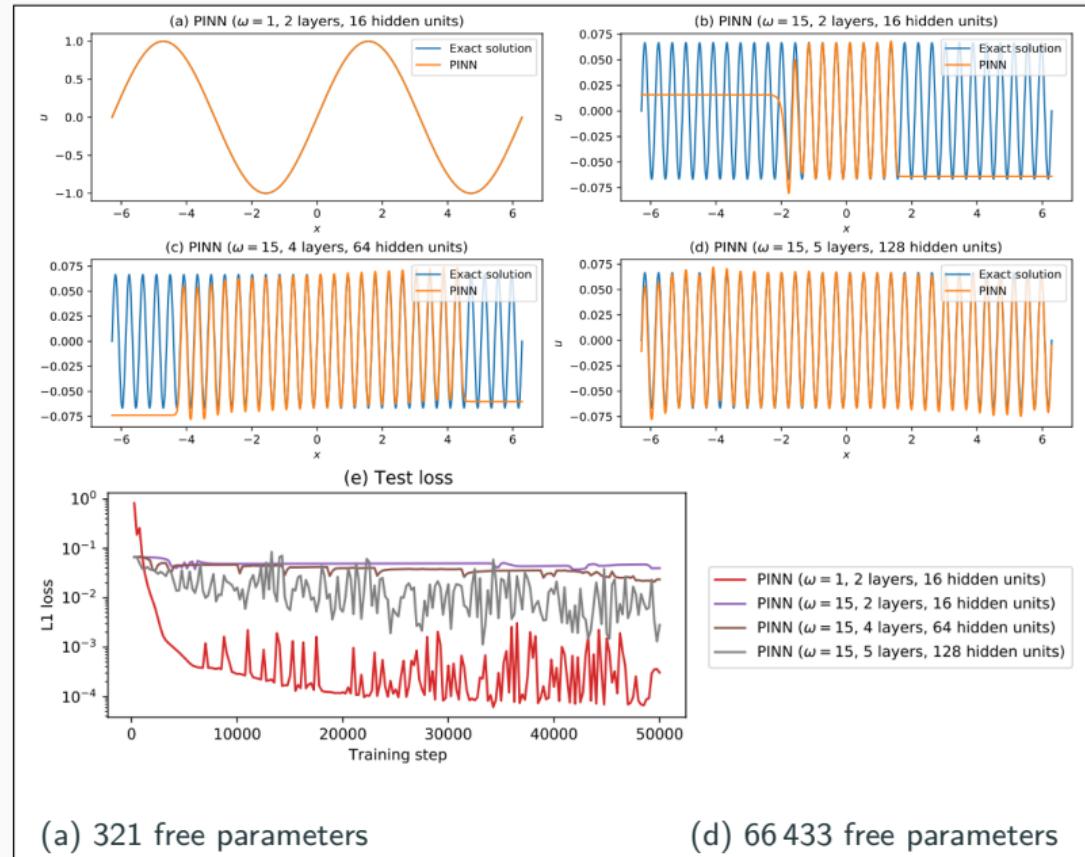
$$\begin{aligned} u' &= \cos(\omega x), \\ u(0) &= 0, \end{aligned}$$

for different values of ω
using PINNs with
varying network
capacities.

Scaling issues

- Large computational domains
- Small frequencies

Cf. Moseley, Markham, and
Nissen-Meyer (2023)



Motivation – Some Observations on the Performance of PINNs

Solve

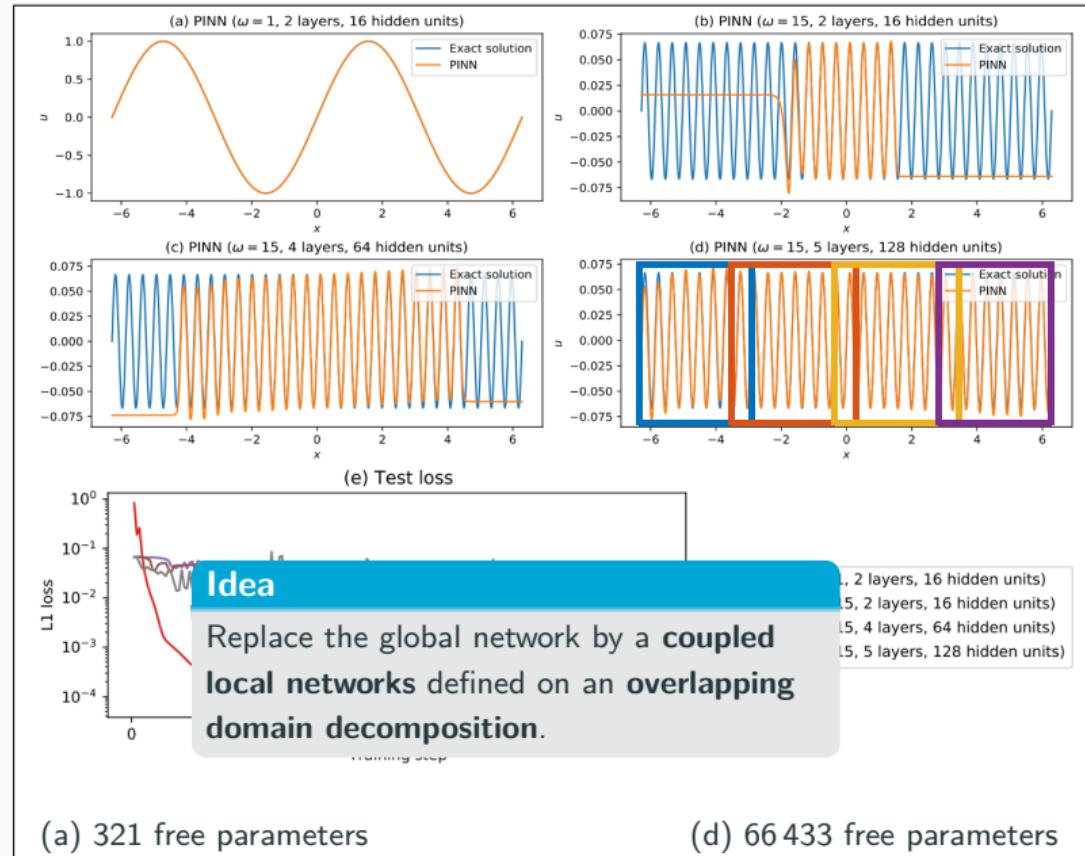
$$\begin{aligned} u' &= \cos(\omega x), \\ u(0) &= 0, \end{aligned}$$

for different values of ω
using PINNs with
varying network
capacities.

Scaling issues

- Large computational domains
- Small frequencies

Cf. Moseley, Markham, and
Nissen-Meyer (2023)



Finite Basis Physics-Informed Neural Networks (FBPINNs)

In the **finite basis physics informed neural network (FBPINNs) method** introduced in **Moseley, Markham, and Nissen-Meyer (2023)**, we employ the **PINN** approach and **hard enforcement of the boundary conditions**; cf. **Lagaris et al. (1998)**.

FBPINNs use the **network architecture**

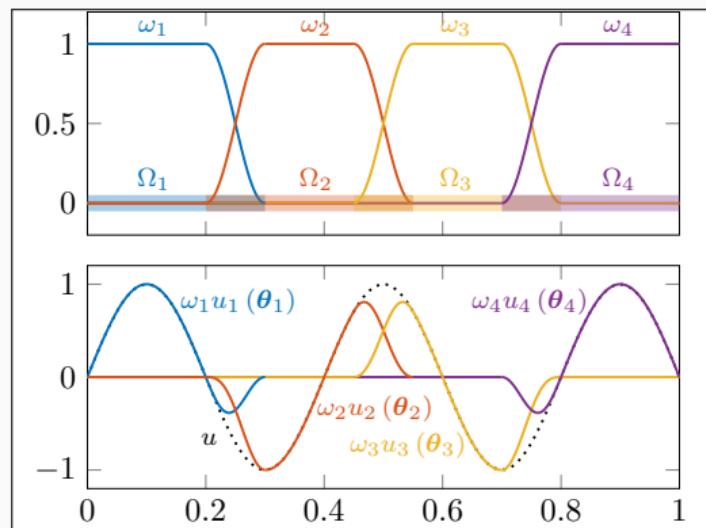
$$u(\theta_1, \dots, \theta_J) = \mathcal{C} \sum_{j=1}^J \omega_j u_j(\theta_j)$$

and the **loss function**

$$\mathcal{L}(\theta_1, \dots, \theta_J) = \frac{1}{N} \sum_{i=1}^N \left(n[\mathcal{C} \sum_{x_i \in \Omega_j} \omega_j u_j](x_i, \theta_j) - f(x_i) \right)^2.$$

Here:

- **Overlapping DD:** $\Omega = \bigcup_{j=1}^J \Omega_j$
- **Partition of unity** ω_j with $\text{supp}(\omega_j) \subset \Omega_j$ and $\sum_{j=1}^J \omega_j \equiv 1$ on Ω



Hard enf. of boundary conditions

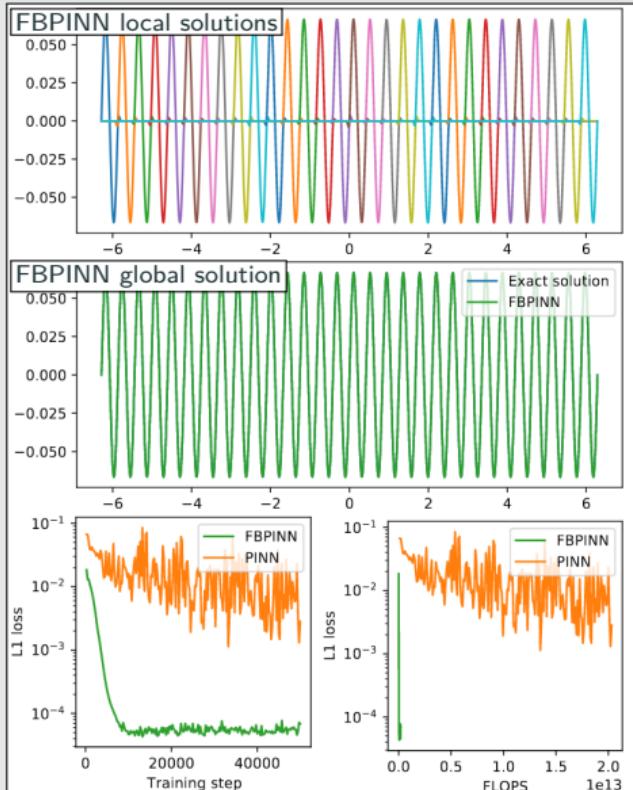
Loss function

$$\mathcal{L}(\theta) = \frac{1}{N} \sum_{i=1}^N (n[\mathcal{C} u](x_i, \theta) - f(x_i))^2,$$

with constraining operator \mathcal{C} , which **explicitly enforces the boundary conditions**.

Numerical Results for FBPINNs

PINN vs FBPINN (Moseley et al. (2023))



Scalability of FBPINNs

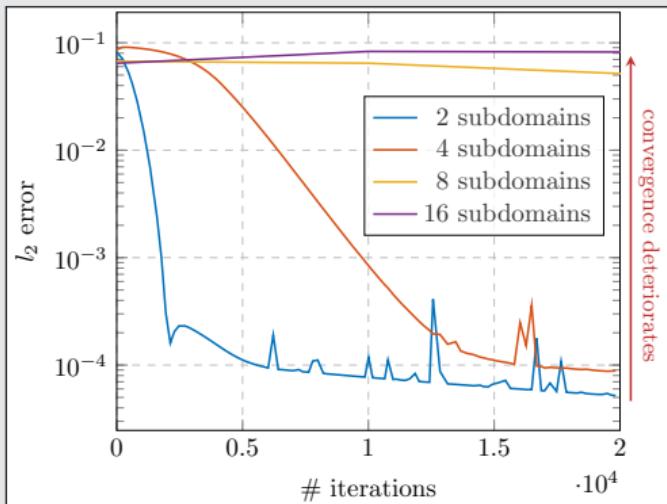
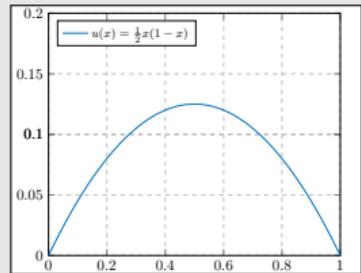
Consider the simple boundary value problem

$$-u'' = 1 \text{ in } [0, 1],$$

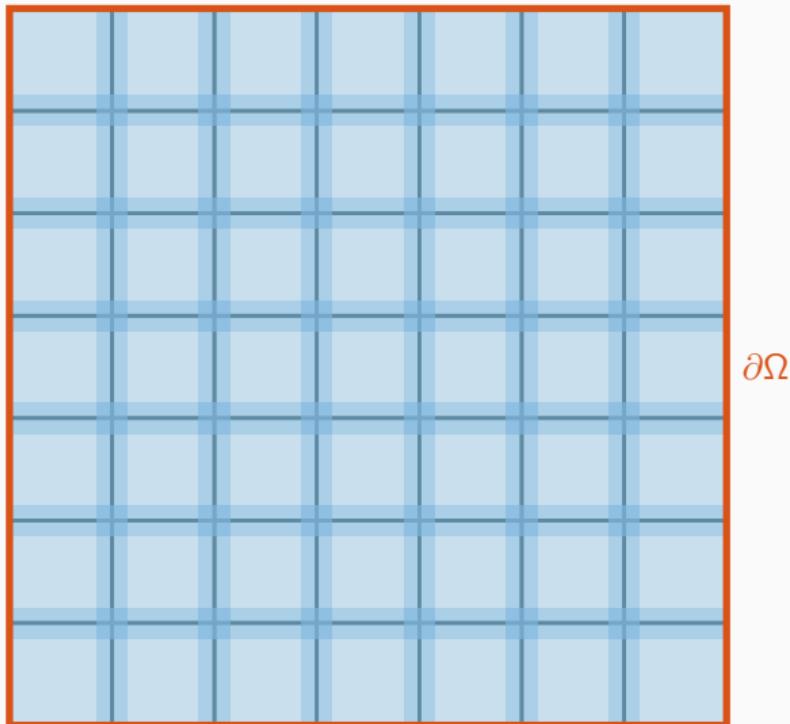
$$u(0) = u(1) = 0,$$

which has the solution

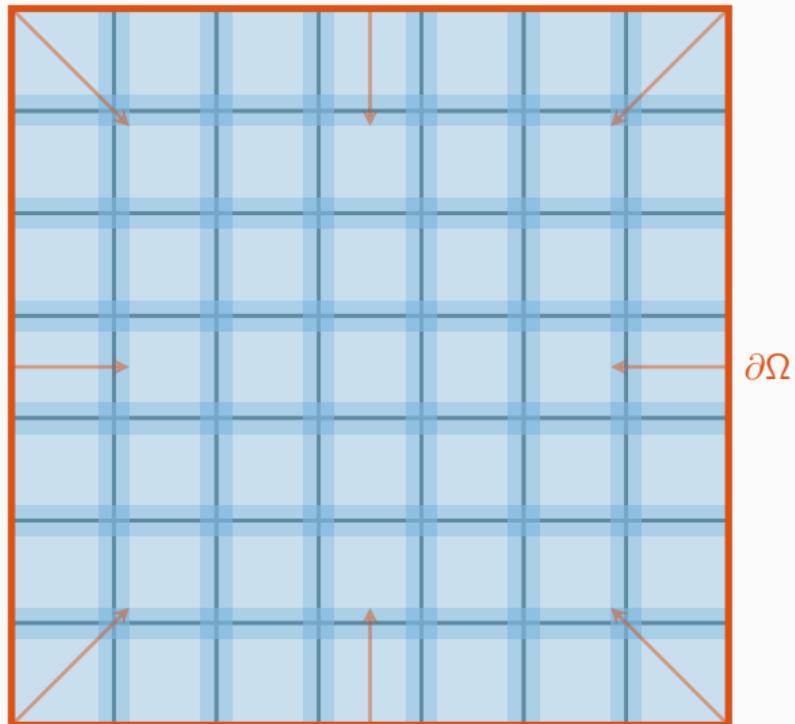
$$u(x) = 1/2x(1 - x).$$



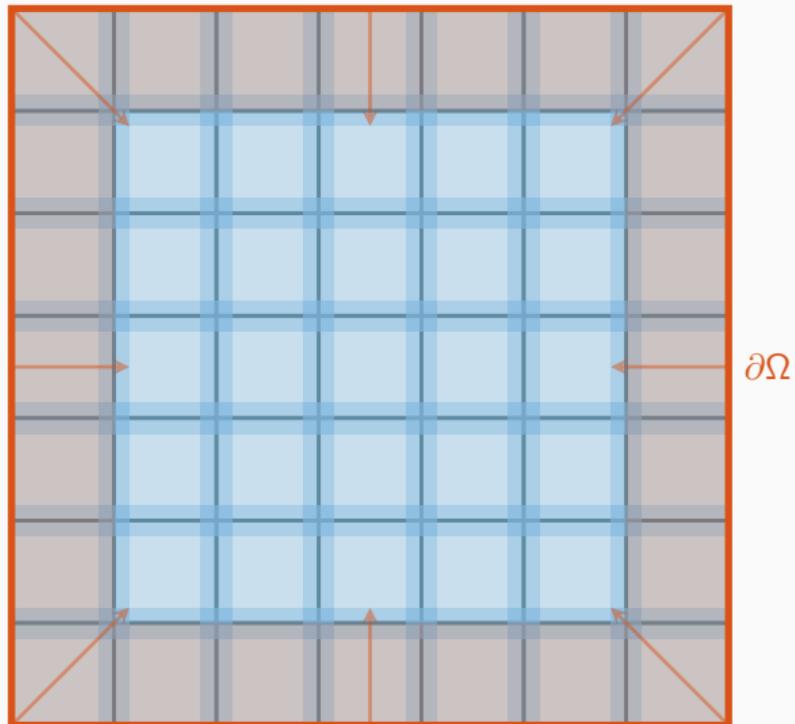
Transport of Information – One-Level Overlapping Schwarz Methods



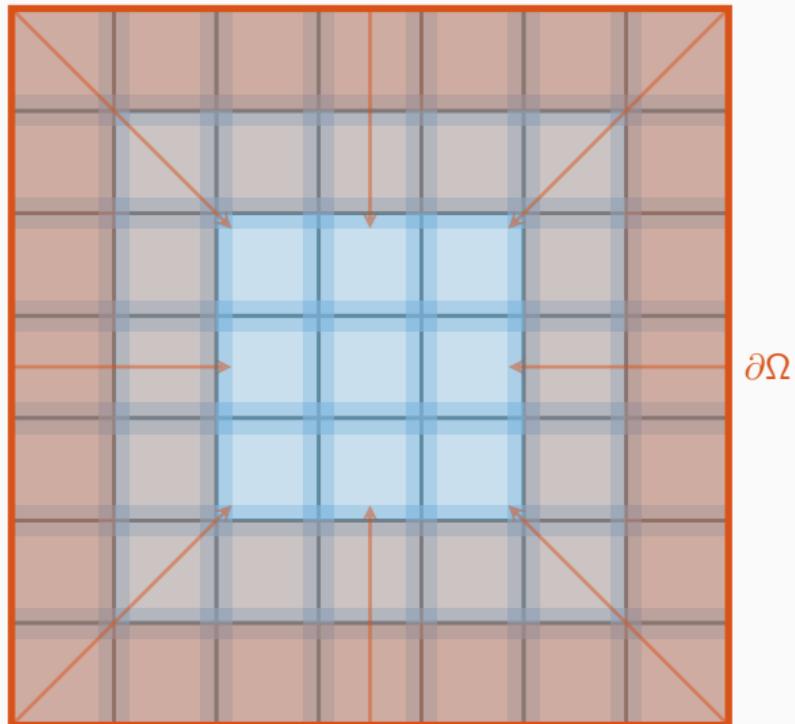
Transport of Information – One-Level Overlapping Schwarz Methods



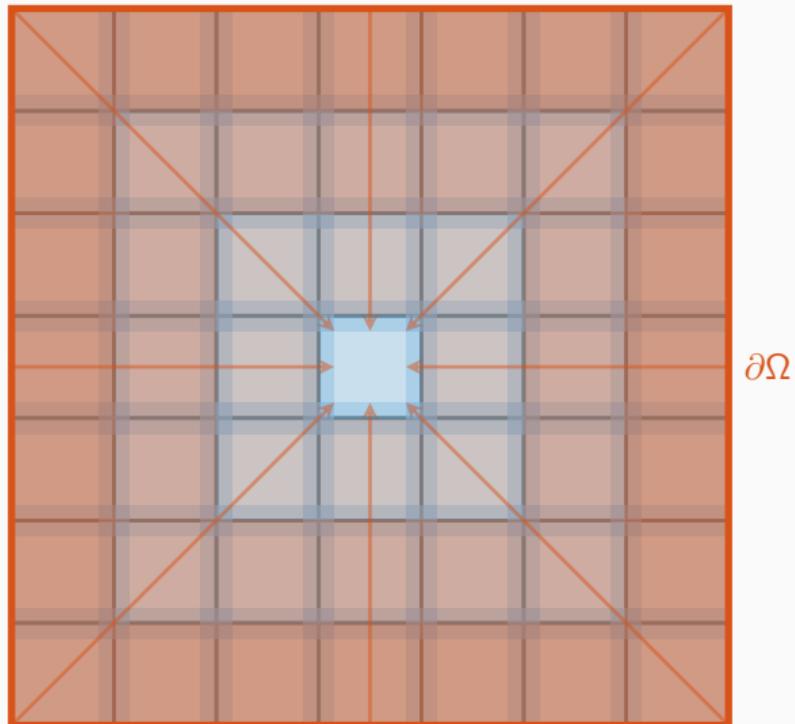
Transport of Information – One-Level Overlapping Schwarz Methods



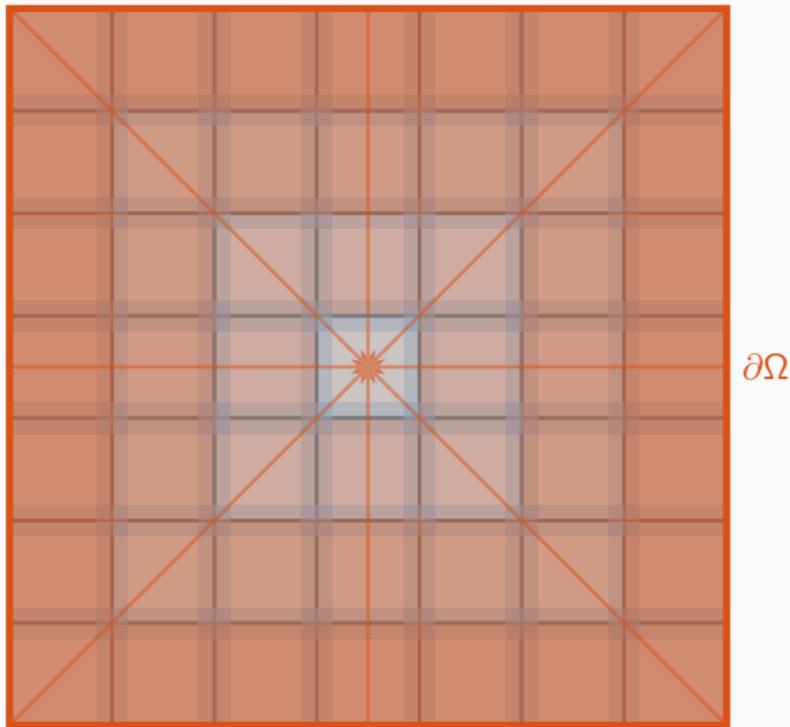
Transport of Information – One-Level Overlapping Schwarz Methods



Transport of Information – One-Level Overlapping Schwarz Methods



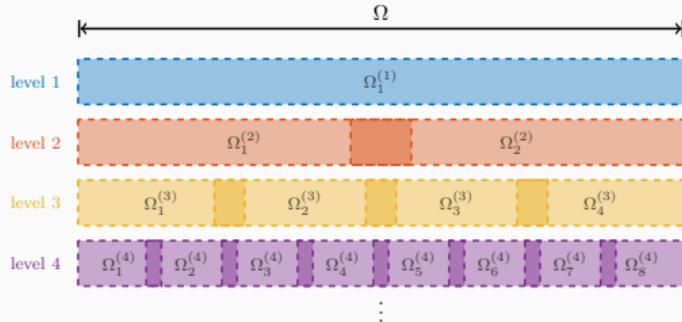
Transport of Information – One-Level Overlapping Schwarz Methods



Information (in particular, boundary data) is **only exchanged via the overlapping regions**, leading to **slow convergence** → establish a faster / global transport of information.

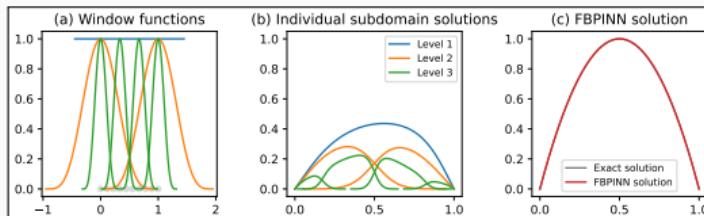
Multi-Level FBPINN Algorithm

Extension of FBPINNs to L levels; Cf. **Dolean, Heinlein, Mishra, Moseley (2024)**.



L -level network architecture

$$u(\theta_1^{(1)}, \dots, \theta_{J(L)}^{(L)}) = \mathcal{C} \left(\sum_{l=1}^L \sum_{i=1}^{N^{(l)}} \omega_j^{(l)} u_j^{(l)}(\theta_j^{(l)}) \right)$$



Multi-Frequency Problem

Let us now consider the two-dimensional multi-frequency Laplace boundary value problem

$$\begin{aligned} -\Delta u &= 2 \sum_{i=1}^n (\omega_i \pi)^2 \sin(\omega_i \pi x) \sin(\omega_i \pi y) && \text{in } \Omega, \\ u &= 0 && \text{on } \partial\Omega, \end{aligned}$$

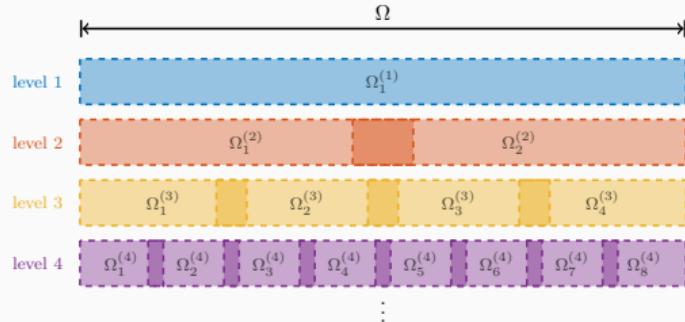
with $\omega_i = 2^i$.

For increasing values of n , we obtain the analytical solutions:



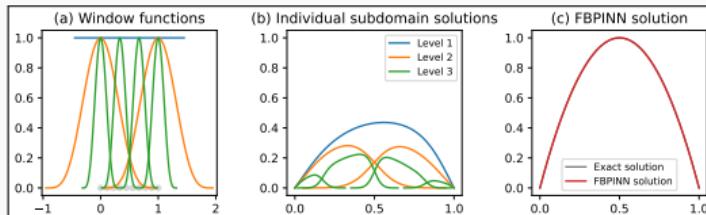
Multi-Level FBPINN Algorithm

Extension of FBPINNs to L levels; Cf. [Dolean, Heinlein, Mishra, Moseley \(2024\)](#).



L -level network architecture

$$u(\theta_1^{(1)}, \dots, \theta_{J(L)}^{(L)}) = \mathcal{C} \left(\sum_{l=1}^L \sum_{i=1}^{N^{(l)}} \omega_j^{(l)} u_j^{(l)}(\theta_j^{(l)}) \right)$$



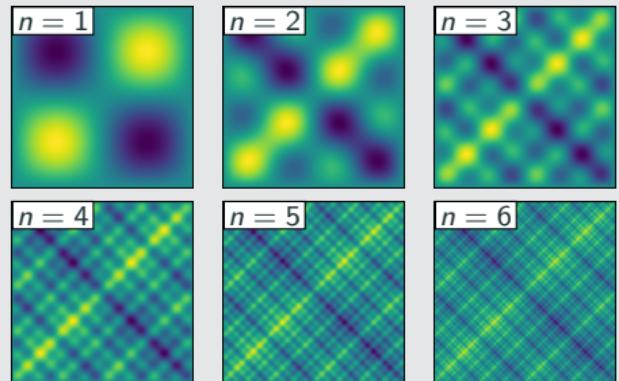
Multi-Frequency Problem

Let us now consider the two-dimensional multi-frequency Laplace boundary value problem

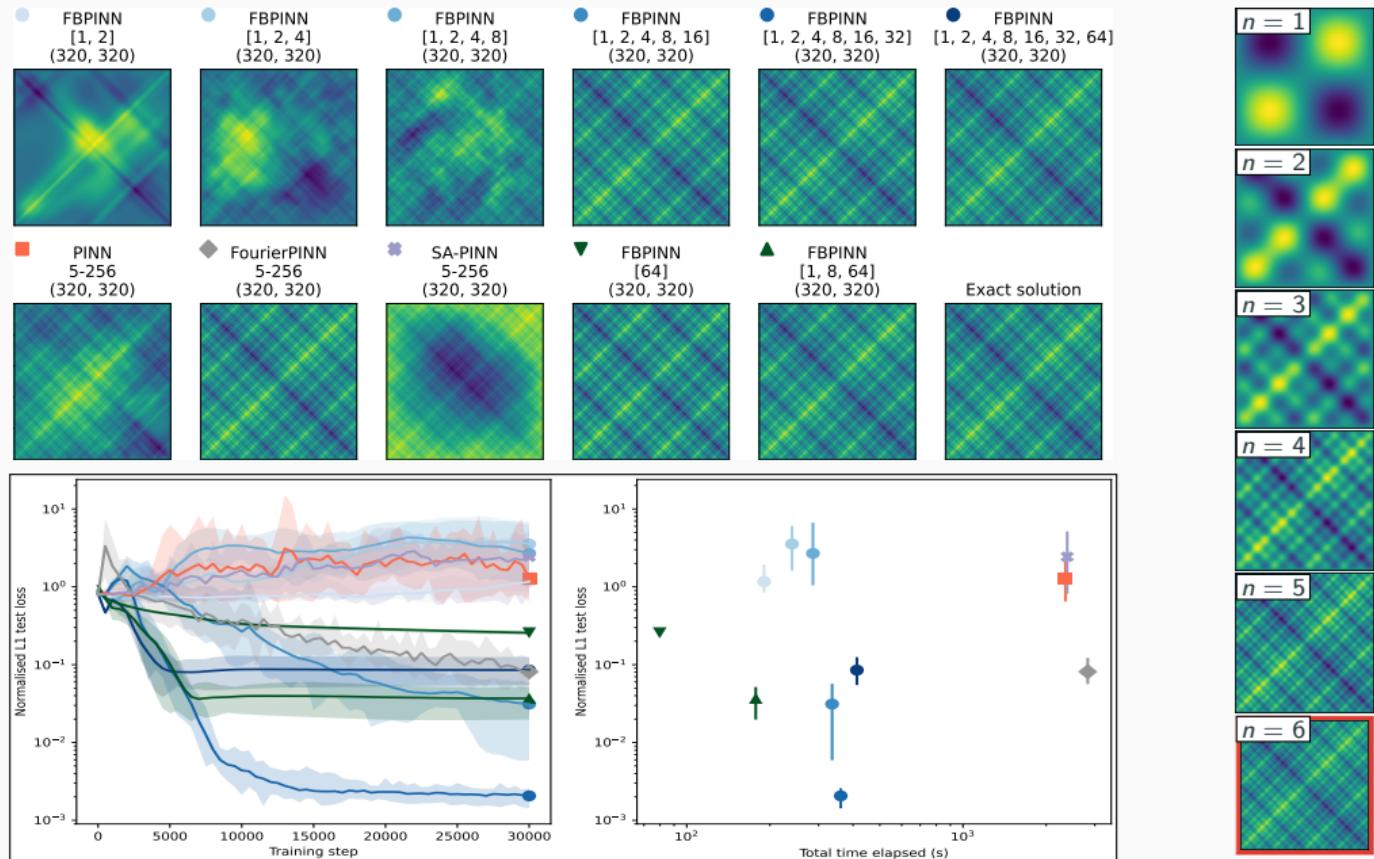
$$-\Delta u = 2 \sum_{i=1}^n (\omega_i \pi)^2 \sin(\omega_i \pi x) \sin(\omega_i \pi y) \quad \text{in } \Omega,$$
$$u = 0 \quad \text{on } \partial\Omega,$$

with $\omega_i = 2^i$.

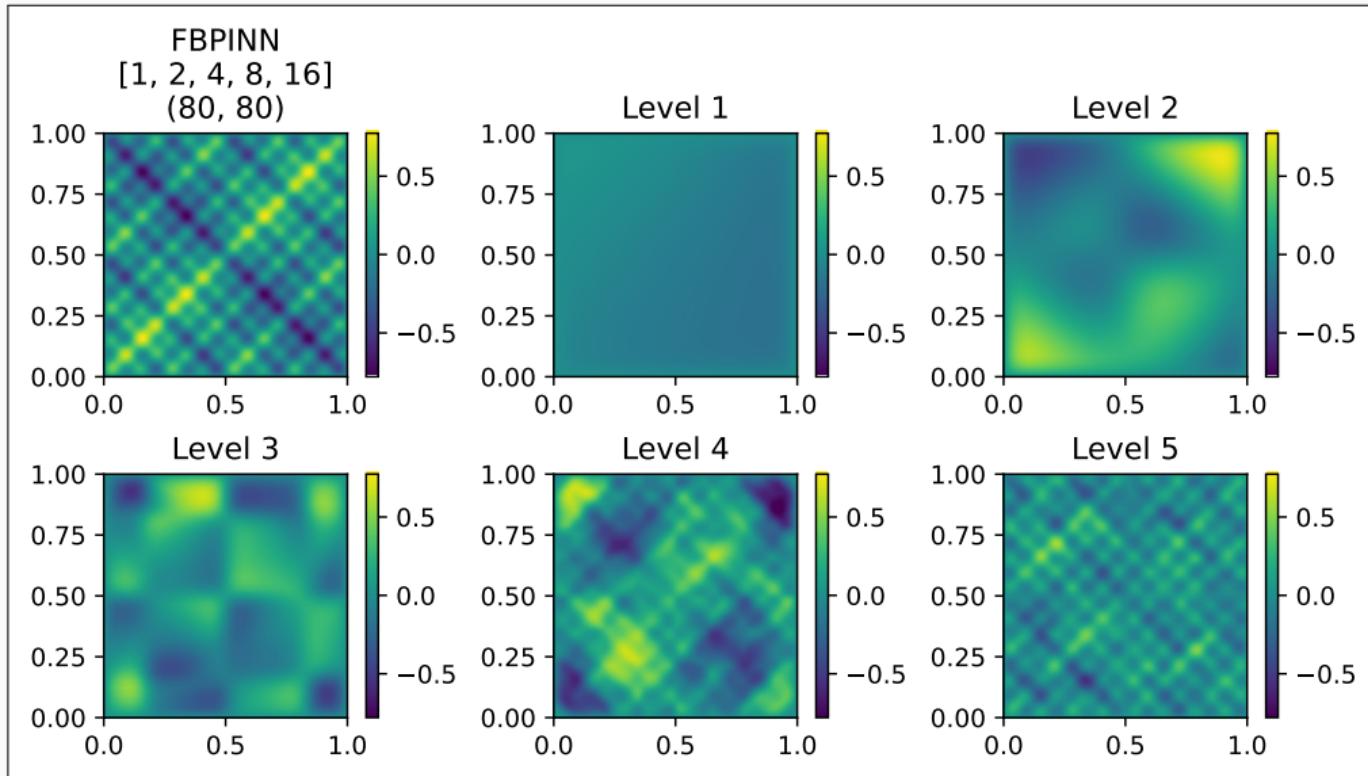
For increasing values of n , we obtain the analytical solutions:



Multi-Level FBPINNs for a Multi-Frequency Problem – Strong Scaling

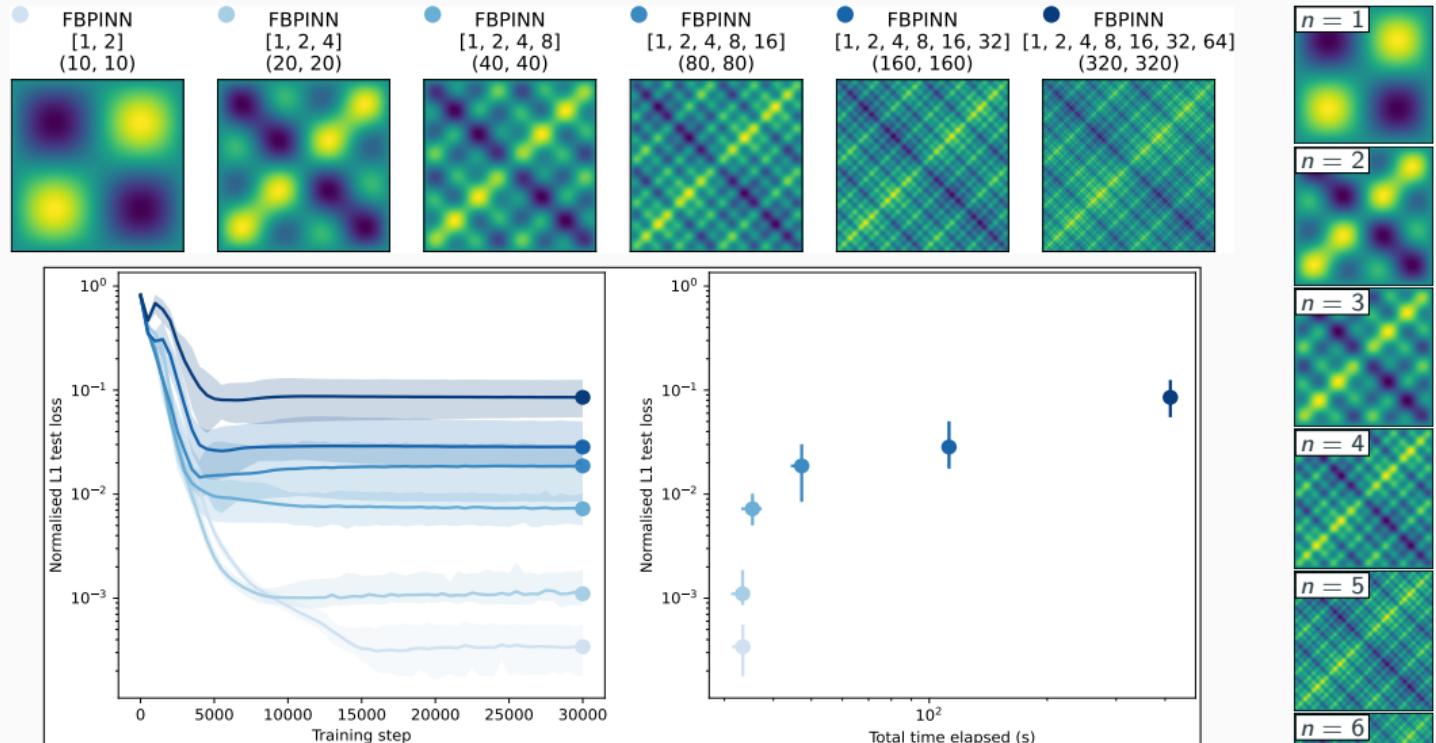


Multi-Frequency Problem – What the FBPINN Learns



Cf. Dolean, Heinlein, Mishra, Moseley (2024).

Multi-Level FBPINNs for a Multi-Frequency Problem – Weak Scaling

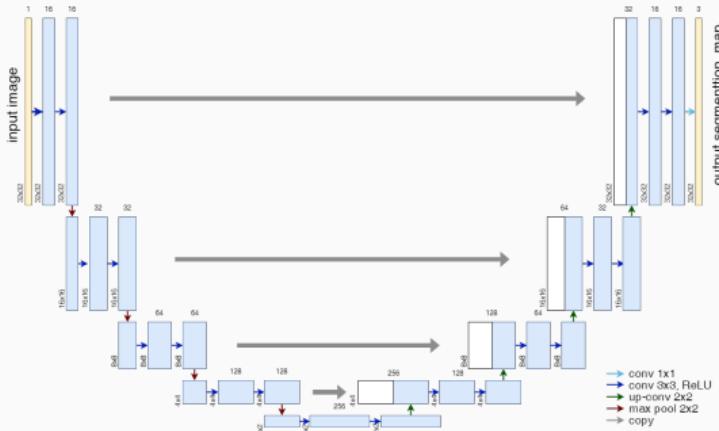


- Ongoing: analysis and improvement of the convergence

Cf. Dolean, Heinlein, Mishra, Moseley (2024).

Domain Decomposition for Convolutional Neural Networks

Memory Requirements for CNN Training

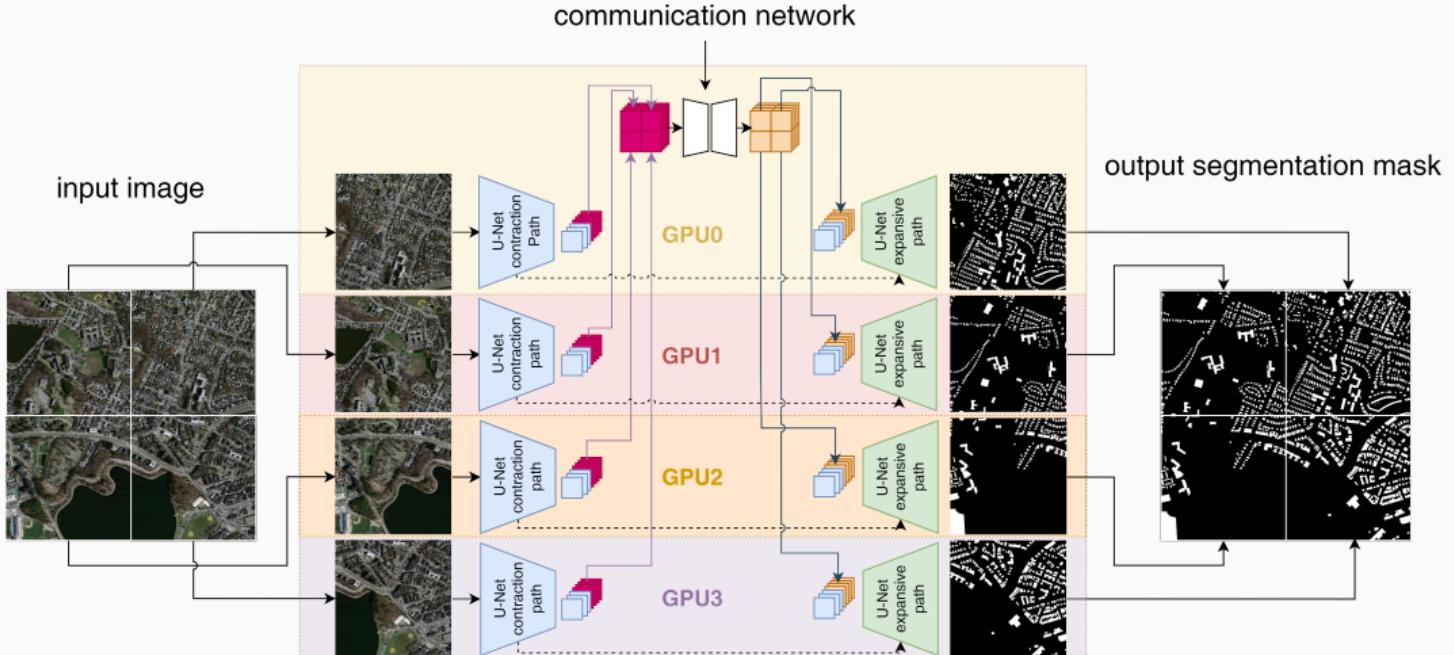


- As an example for a **convolutional neural network (CNN)**, we employ the **U-Net architecture** introduced in **Ronneberger, Fischer, and Brox (2015)**.
- The U-Net yields **state-of-the-art accuracy in semantic image segmentation** and other **image-to-image tasks**.

Below: memory consumption for training on a single 1024×1024 image.

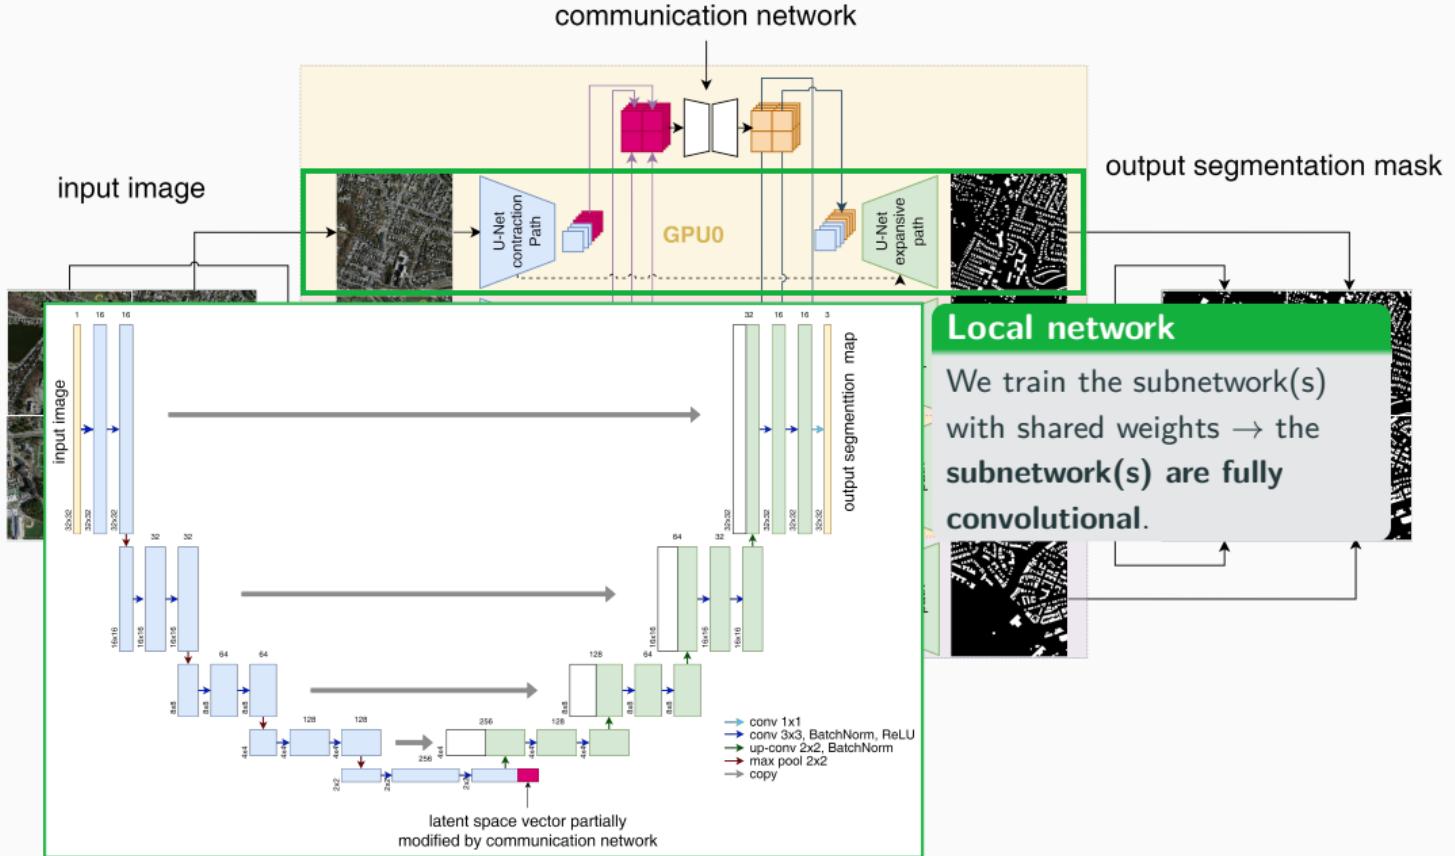
name	size	# channels		mem. feature maps		mem. weights	
		input	output	# of values	MB	# of values	MB
input block	1 024	3	64	268 M	1 024.0	38 848	0.148
encoder block 1	512	64	128	167 M	704.0	221 696	0.846
encoder block 2	256	128	256	84 M	352.0	885 760	3.379
encoder block 3	128	256	512	42 M	176.0	3 540 992	13.508
encoder block 4	64	512	1 024	21 M	88.0	14 159 872	54.016
decoder block 1	64	1,024	512	50 M	192.0	9 177 088	35.008
decoder block 2	128	512	256	101 M	384.0	2 294 784	8.754
decoder block 3	256	256	128	201 M	768.0	573 952	2.189
decoder block 4	512	128	64	402 M	1 536.0	143 616	0.548
output block	1 024	64	3	3.1 M	12.0	195	0.001

Decomposing the U-Net

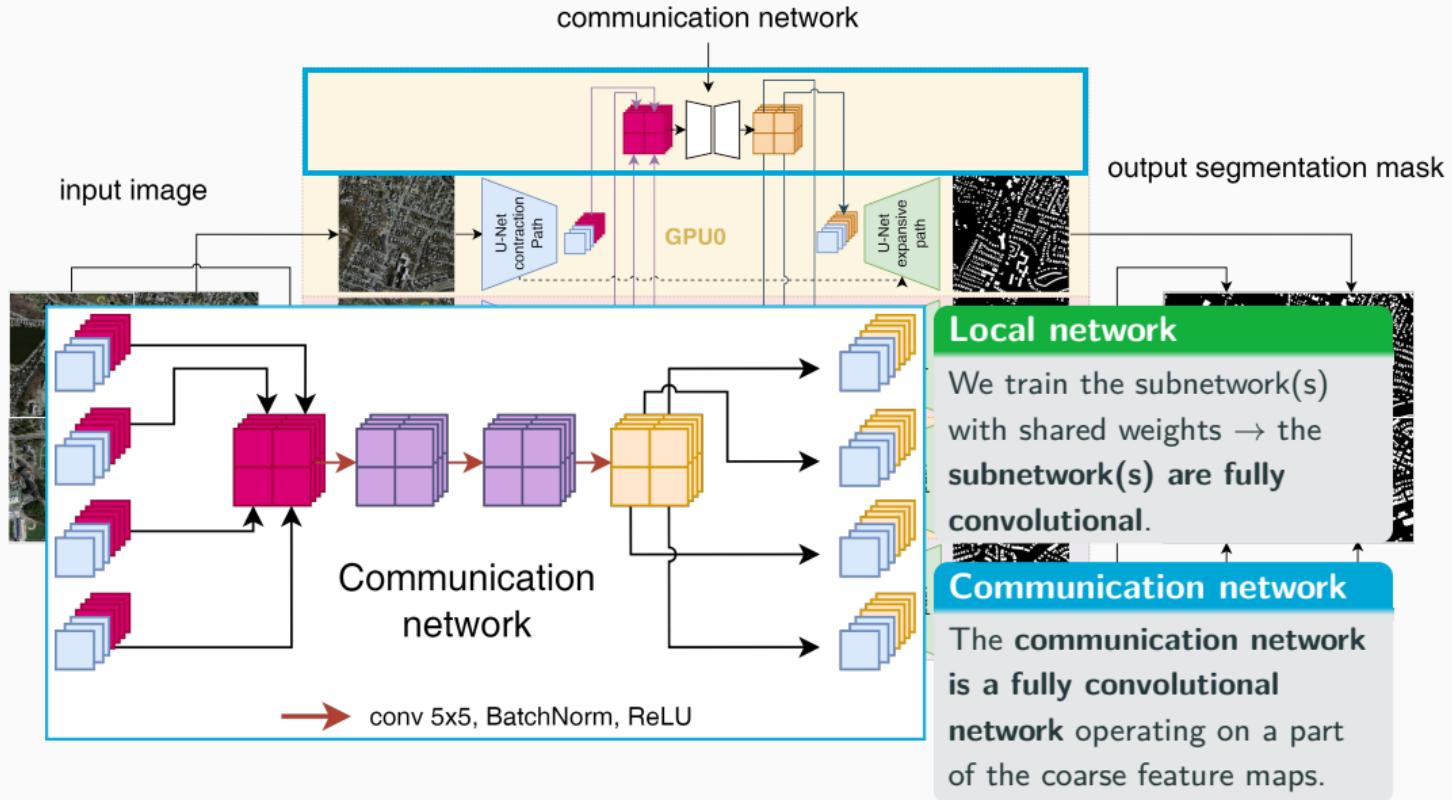


Cf. Verburg, Heinlein, Cyr (subm. 2024).

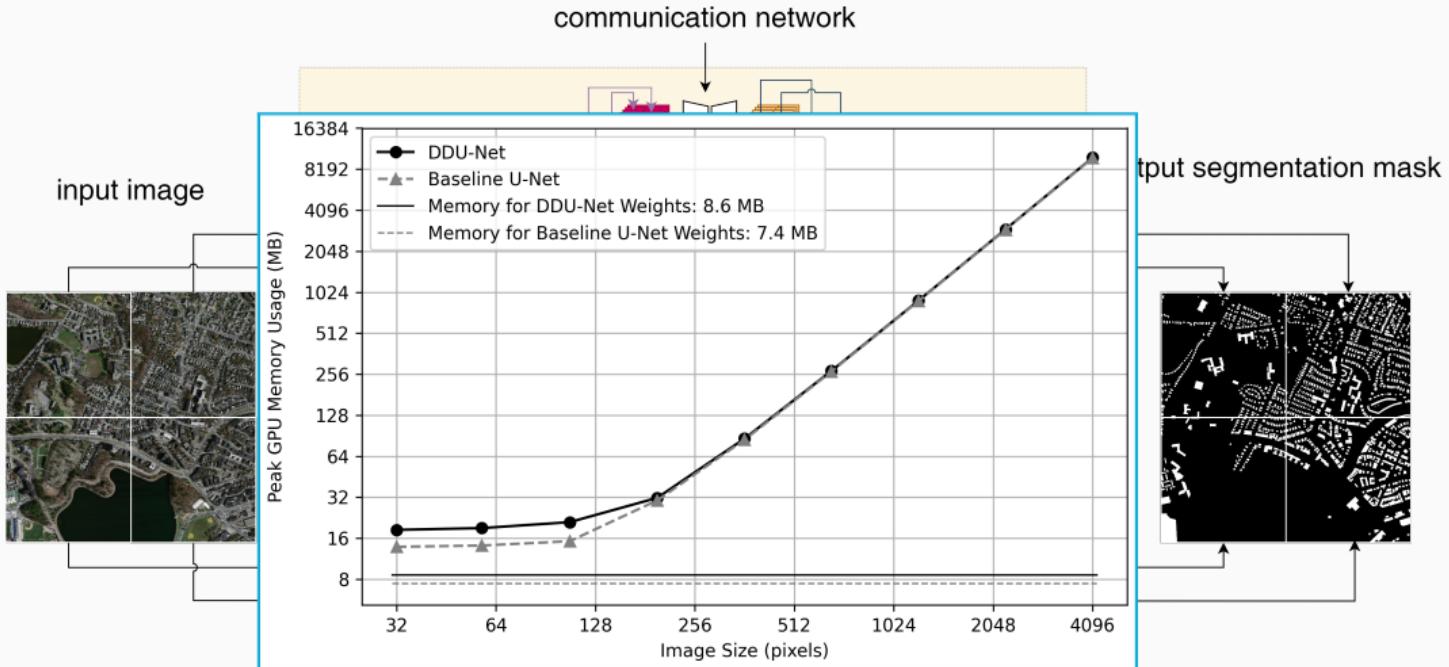
Decomposing the U-Net



Decomposing the U-Net



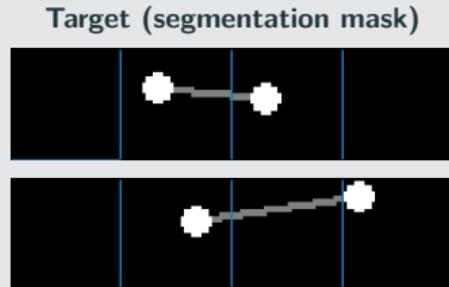
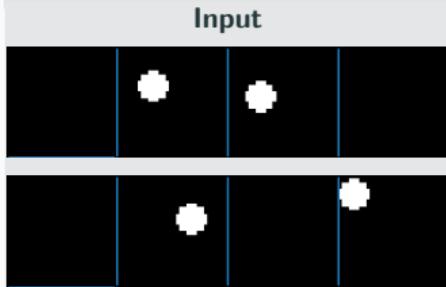
Decomposing the U-Net



- Distribution of feature maps results in **significant reduction of memory usage on a single GPU**
- Moderate **additional memory usage** due to the **communication network**

Results – Synthetic Data Set

Task: Connect two dots via a line segment



Result: Communication

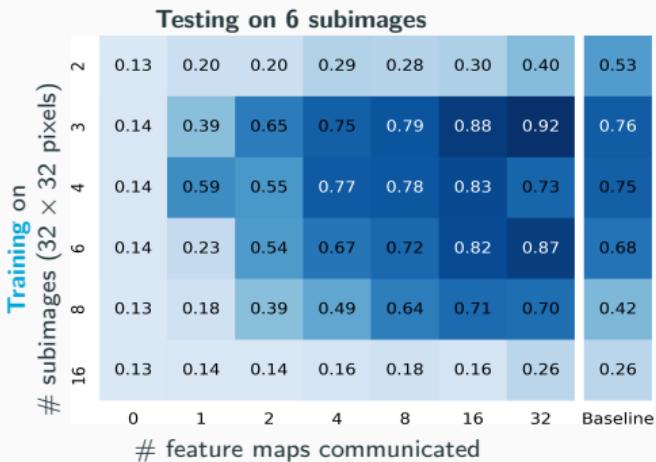
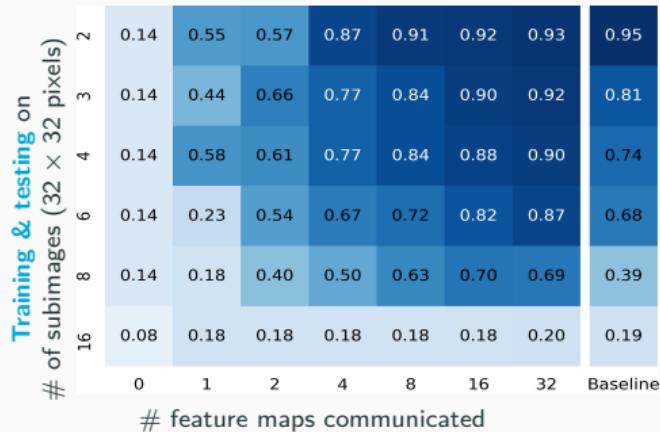
True mask



Pred. (no comm.)

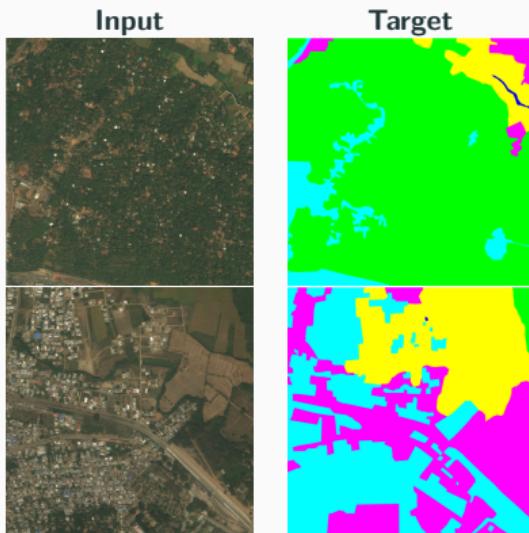


Pred. (comm.)



DeepGlobe 2018 Satellite Image Data Set (Demir et al. (2018))

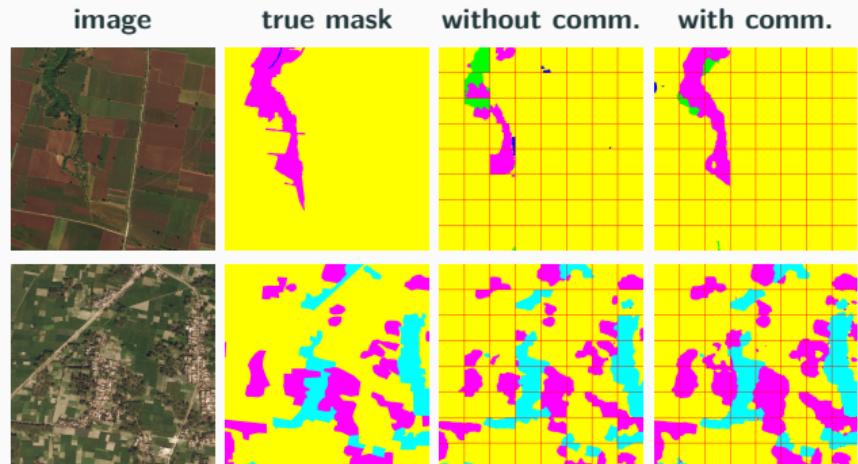
class	pixel count	proportion
urban	642.4M	9.35 %
agriculture	3898.0M	56.76 %
rangeland	701.1M	10.21 %
forest	944.4M	13.75 %
water	256.9M	3.74 %
barren	421.8M	6.14 %
unknown	3.0M	0.04 %



Avoiding overfitting

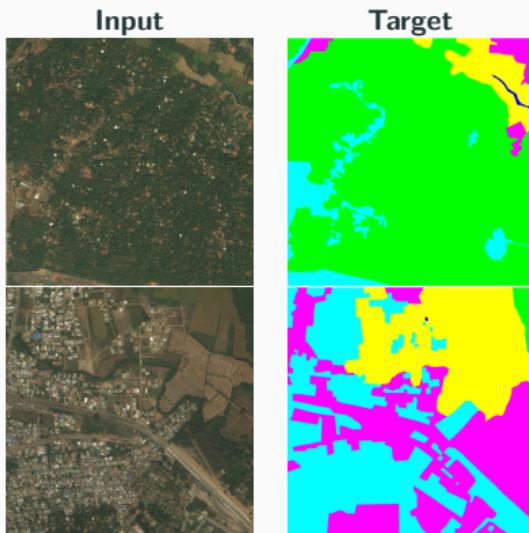
The data set includes **only 803 images**. To **avoid overfitting**, we

- apply **batch normalization**, use **random dropout** layers and **data augmentation**, and
- initialize the encoder using the **ResNet-18** (He, Zhang, Ren, and Sun (2016))



DeepGlobe 2018 Satellite Image Data Set (Demir et al. (2018))

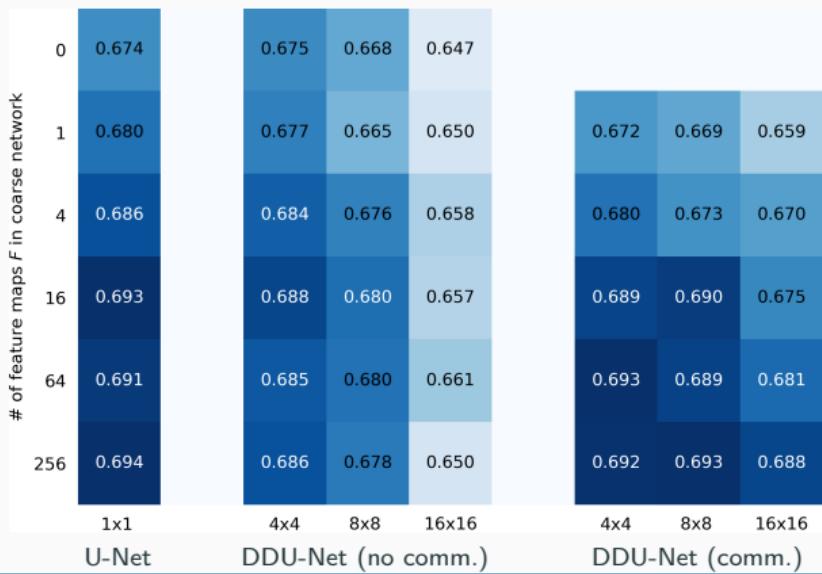
class	pixel count	proportion
urban	642.4M	9.35 %
agriculture	3898.0M	56.76 %
rangeland	701.1M	10.21 %
forest	944.4M	13.75 %
water	256.9M	3.74 %
barren	421.8M	6.14 %
unknown	3.0M	0.04 %



Avoiding overfitting

The data set includes **only 803 images**. To **avoid overfitting**, we

- apply **batch normalization**, use **random dropout** layers and **data augmentation**, and
- initialize the encoder using the **ResNet-18** (He, Zhang, Ren, and Sun (2016))



FBPINNs – Domain Decomposition for Physics-Informed Neural Networks

- Schwarz domain decomposition architectures **improve the scalability of PINNs** to **large domains / high frequencies**, **keeping the complexity of the local networks low**.
- As classical domain decomposition methods, **one-level FBPINNs** are **not scalable to large numbers of subdomains**; **multilevel FBPINNs enable scalability**.

DDU-Net – Domain Decomposition for CNNs

- The **memory requirements for training of high-resolution images** using CNNs can be **large**. In particular, the U-Net model requires **storing intermediate feature maps**.
- Our **novel DDU-Net** approach **decouples the training on the sub-images**, allowing us to **distribute the memory load** among **multiple GPUs**. It **limits communication** to **deepest level** of the U-Net architecture using a **communication network**.

Thank you for your attention!