






Localization of Neural Networks Using Domain Decomposition

Alexander Heinlein¹

Seminar, Great Bay University, Dongguan, China, November 10, 2025

¹Delft University of Technology

Scientific Computing and Machine Learning



Numerical methods

Based on physical models

- + Robust and generalizable
- Require availability of mathematical models

Machine learning models

Driven by data

- + Do not require mathematical models
- Sensitive to data, limited extrapolation capabilities

Scientific machine learning

Combining the strengths and compensating the weaknesses of the individual approaches:

numerical methods	improve	machine learning techniques
machine learning techniques	assist	numerical methods

Outline

1 Domain decomposition for physics-informed neural networks

Based on joint work with

Victorita Dolean

(Eindhoven University of Technology)

Taniya Kapoor

(Wageningen University & Research)

Siddhartha Mishra

(ETH Zürich)

Ben Moseley

(Imperial College London)

Yong Shang and Fei Wang

(Xi'an Jiaotong University)

2 Spectral analysis and domain decomposition for deep operator networks

Based on joint work with

Amanda A. Howard and Panos Stinis

(Pacific Northwest National Laboratory)

Johannes Taraz

(Delft University of Technology)

3 Domain decomposition-based image segmentation for high-resolution image segmentation on multiple GPUs

Based on joint work with

Eric Cyr

(Sandia National Laboratories)

Corné Verburg

(Delft University of Technology)

Domain decomposition for physics-informed neural networks

Physics-Informed Neural Networks (PINNs) – Idea


In [Lagaris et al. \(1998\)](#), the authors solve the boundary value problem

$$-\Delta \Psi_t(x, \theta) = 1 \text{ on } [0, 1],$$

$$\Psi_t(0, \theta) = \Psi_t(1, \theta) = 0,$$

via a **collocation approach**:

$$\min_{\theta} \sum_{x_i} (\Delta \Psi_t(x_i, \theta) + 1)^2$$




Boundary conditions ...

... can be enforced explicitly via the ansatz:

$$\Psi_t(x, \theta) = A(x) + F(x, \text{NN}(x, \theta))$$

- A satisfies the boundary conditions
- F does not contribute to the boundary conditions



Physics-Informed Neural Networks (PINNs)

In the **physics-informed neural network (PINN)** approach introduced by **Raissi et al. (2019)**, a **neural network** is employed to **discretize a partial differential equation**

$$\mathcal{N}[u] = f, \quad \text{in } \Omega.$$

PINNs use a **hybrid loss function**:



$$\mathcal{L}(\theta) = \omega_{\text{data}} \mathcal{L}_{\text{data}}(\theta) + \omega_{\text{PDE}} \mathcal{L}_{\text{PDE}}(\theta),$$

where ω_{data} and ω_{PDE} are **weights** and


$$\mathcal{L}_{\text{data}}(\theta) = \frac{1}{N_{\text{data}}} \sum_{i=1}^{N_{\text{data}}} (u(\hat{x}_i, \theta) - u_i)^2,$$

$$\mathcal{L}_{\text{PDE}}(\theta) = \frac{1}{N_{\text{PDE}}} \sum_{i=1}^{N_{\text{PDE}}} (\mathcal{N}[u](x_i, \theta) - f(x_i))^2.$$

See also Dissanayake and Phan-Thien (1994); Lagaris et al. (1998).



Hybrid loss



Advantages

- "Meshfree"
- Small data
- Generalization properties
- High-dimensional problems
- Inverse and parameterized problems

Drawbacks

- Training cost and robustness
- Convergence not well-understood
- Difficulties with scalability and multi-scale problems

- Known solution values can be included in $\mathcal{L}_{\text{data}}$
- Initial and boundary conditions are also included in $\mathcal{L}_{\text{data}}$

Error Estimate & Spectral Bias


Estimate of the generalization error ([Mishra and Molinaro \(2022\)](#))

The generalization error (or total error) satisfies

$$\mathcal{E}_G \leq C_{\text{PDE}} \mathcal{E}_{\mathcal{T}} + C_{\text{PDE}} C_{\text{quad}}^{1/p} N^{-\alpha/p}$$

- $\mathcal{E}_G = \mathcal{E}_G(\mathbf{X}, \theta) := \|\mathbf{u} - \mathbf{u}^*\|_V$ **general. error** (V Sobolev space, \mathbf{X} training data set)
- $\mathcal{E}_{\mathcal{T}}$ **training error** (l^p loss of the residual of the PDE)
- N **number of the training points** and α **convergence rate of the quadrature**
- C_{PDE} and C_{quad} **constants** depending on the **PDE, quadrature, and neural network**

Rule of thumb: “As long as the PINN is **trained well**, it also **generalizes well**”




[Rahaman et al., On the spectral bias of neural networks, ICML \(2019\)](#)

Related works: [Cao et al. \(2021\)](#), [Wang, et al. \(2022\)](#), [Hong et al. \(arXiv 2022\)](#), [Xu et al. \(2024\)](#), ...

Spectral Bias of Neural Networks

Rahaman et al. (2019) observed the **spectral bias of neural networks**: during training, they learn low-frequency functions faster than high-frequency functions; see also Cao et al. (2021), Wang, et al. (2022), Hong et al. (arXiv 2022), Xu et al. (2024), ...



This can be understood by interpreting the training dynamics of neural networks as **gradient flow**

$$\theta_{k+1} = \theta_k - \eta_k \nabla_\theta \mathcal{C}(n_{\theta_k}) \quad \rightarrow \quad \frac{dx}{dt}(t) = -\nabla_x \mathcal{C}(x(t))$$


For **infinitely wide neural networks**, we obtain a **constant neural tangent kernel (NTK)** (Jacot et al. (2018)).

For mean squared error (MSE) loss \mathcal{C} , we obtain the **discretized K** = $\left(\left\langle \frac{d n_{\theta(t)}}{d \theta}(x_i), \frac{d n_{\theta(t)}}{d \theta}(x_j) \right\rangle \right)_{ij}$:

$$\frac{d n_{\theta(t)}}{d t} = -\frac{2}{n} \cdot K(t) (n_{\theta(t)}(\mathbf{X}) - \mathbf{Y}) \quad \Rightarrow \quad \mathbf{U}^\top (n_{\theta(t)}(\mathbf{X}) - \mathbf{Y}) \approx e^{-t \frac{2}{n} \Lambda} \mathbf{U}^\top \mathbf{Y}$$

Spectral Bias of Neural Networks

Rahaman et al. (2019) observed the **spectral bias of neural networks**: during training, they learn **low-frequency functions faster than high-frequency functions**; see also Cao et al. (2021), Wang, et al. (2022), Hong et al. (arXiv 2022), Xu et al. (2024), ...



Scaling of PINNs for a Simple ODE Problem

Solve


$$\begin{aligned} u' &= \cos(\omega x), \\ u(0) &= 0, \end{aligned}$$

for different values of ω
using PINNs with
varying network
capacities.

Scaling issues

- Large computational domains
- Small frequencies

Cf. Moseley, Markham, and
Nissen-Meyer (2023)



Scaling of PINNs for a Simple ODE Problem

Solve


$$\begin{aligned} u' &= \cos(\omega x), \\ u(0) &= 0, \end{aligned}$$

for different values of ω
using PINNs with
varying network
capacities.


Scaling issues

- Large computational domains
- Small frequencies

Cf. Moseley, Markham, and Nissen-Meyer (2023)




Domain Decomposition Methods



Graphics based on results from Heinlein, Perego, Rajamanickam (2022)

Historical remarks: The **alternating Schwarz method** is the earliest **domain decomposition method (DDM)**, which has been invented by **H. A. Schwarz** and published in **1870**:

- Schwarz used the algorithm to establish the **existence of harmonic functions** with prescribed boundary values on **regions with non-smooth boundaries**.



Domain Decomposition Methods and Machine Learning – Literature

A non-exhaustive literature overview:

- ML for adaptive BDDC, FETI–DP, AGDSW: H., Klawonn, Langer, Weber (2019, 2020, 2021, 2021, 2021, 2022); Klawonn, Langer, Weber (2024, 2025)
- cPINNs, XPINNs: Jagtap, Kharazmi, Karniadakis (2020); Jagtap, Karniadakis (2020)
- Classical Schwarz iteration for PINNs or DeepRitz:: Li, Tang, Wu, and Liao (2019); Li, Xiang, Xu (2020); Mercier, Gratton, Boudier (arXiv 2021); Dolean, H., Mercier, Gratton (acc. 2025); Li, Wang, Cui, Xiang, Xu (2023); Sun, Xu, Yi (arXiv 2023, 2024); Kim, Yang (2023, 2024, 2024)
- FBPINNs, FBKANs: Moseley, Markham, Nissen-Meyer (2023); Dolean, H., Mishra, Moseley (2024, 2024); H., Howard, Beecroft, Stinis (2025); Howard, Jacob, Murphy, H., Stinis (arXiv 2024)
- DD for randomized NNs: Dong, Li (2021); Dang, Wang (2024); Sun, Dong, Wang (2024); Sun, Wang (2024); Chen, Chi, E, Yang (2022); Shang, H., Mishra, Wang (2025); Anderson, Dolean, Moseley, Pestana, (arXiv 2024); van Beek, Dolean, Moseley (arxiv 2025)
- DD for Neural Operators and Surrogate Models: H., Howard, Beecroft, Stinis (2025); Ramezankhani, Parekh, Deodhar, Birru (arXiv 2024); Wu, Kovachki, Liu (arXiv 2025); Pelzer, Verburg, H., Schulte (arXiv 2025); Klaes, Klawonn, Kubicki, Langer, Nakajima, Shimokawabe, Weber (arXiv 2025); Howard, H., Stinis (in prep.)
- DD for CNNs: Gu, Zhang, Liu, Cai (2022); Lee, Park, Lee (2022); Klawonn, Langer, Weber (2024); Verburg, H., Cyr (2025)

An overview of the state-of-the-art in 2024:



A. Klawonn, M. Langer, J. Weber

Machine learning, domain decomposition methods – a survey

Computational Science and Engineering. 2024

Finite Basis Physics-Informed Neural Networks (FBPINNs)


FBPINNs ([Moseley, Markham, Nissen-Meyer \(2023\)](#))

FBPINNs employ the **network architecture**


$$u(\theta_1, \dots, \theta_J) = \sum_{j=1}^J \omega_j u_j(\theta_j)$$

and the **loss function**


$$\mathcal{L} = \frac{1}{N} \sum_{i=1}^N \left(n \left[\sum_{x_i \in \Omega_j} \omega_j u_j(x_i, \theta_j) - f(x_i) \right] \right)^2$$



1D single-frequency problem



[Moseley, Markham, Nissen-Meyer \(2023\)](#)



Finite Basis Physics-Informed Neural Networks (FBPINNs)


FBPINNs ([Moseley, Markham, Nissen-Meyer \(2023\)](#))

FBPINNs employ the **network architecture**


$$u(\theta_1, \dots, \theta_J) = \sum_{j=1}^J \omega_j u_j(\theta_j)$$

and the **loss function**

$$\mathcal{L} = \frac{1}{N} \sum_{i=1}^N \left(n \left[\sum_{x_i \in \Omega_j} \omega_j u_j(x_i, \theta_j) - f(x_i) \right] \right)^2$$




1D single-frequency problem



Multi-Level FBPINNs

Multi-level FBPINNs (ML-FBPINNs)

ML-FBPINNs ([Dolean, Heinlein, Mishra, Moseley \(2024\)](#)) are based on a **hierarchy of domain decompositions**:




This yields the **network architecture**:


$$u(\theta_1^{(1)}, \dots, \theta_{J^{(L)}}^{(L)}) = \sum_{l=1}^L \sum_{i=1}^{N^{(l)}} \omega_j^{(l)} u_j^{(l)}(\theta_j^{(l)})$$

Multiscale problems ...


... appear in most areas of modern science and engineering:



Dual-phase steel;
fig. courtesy of
[J. Schröder](#).



Groundwater flow;
cf. [Christie & Blunt \(2001\)](#) (SPE10).



Arterial walls;
cf. [O'Connell et al. \(2008\)](#).

Multi-frequency problem

Consider the **multi-frequency Laplace problem**

$$-\Delta u = 2 \sum_{i=1}^n (\omega_i \pi)^2 \sin(\omega_i \pi x) \sin(\omega_i \pi y),$$

with homogeneous Dirichlet boundary conditions and $\omega_i = 2^i$.

For increasing values of n , we obtain the **solutions**:



$n = 1$



$n = 3$




$n = 5$

Multi-Level FBPINNs

Multi-level FBPINNs (ML-FBPINNs)

ML-FBPINNs ([Dolean, Heinlein, Mishra, Moseley \(2024\)](#)) are based on a **hierarchy** of domain decompositions:



This yields the **network architecture**:


$$u(\theta_1^{(1)}, \dots, \theta_{J^{(L)}}^{(L)}) = \sum_{l=1}^L \sum_{i=1}^{N^{(l)}} \omega_j^{(l)} u_j^{(l)}(\theta_j^{(l)})$$

Multiscale problems ...


... appear in most areas of modern science and engineering:



Dual-phase steel;
fig. courtesy of
[J. Schröder](#).



Groundwater flow;
cf. [Christie & Blunt \(2001\)](#) (SPE10).



Arterial walls;
cf. [O'Connell et al. \(2008\)](#).


Multi-frequency problem

Consider the **multi-frequency Laplace problem**


$$-\Delta u = 2 \sum_{i=1}^n (\omega_i \pi)^2 \sin(\omega_i \pi x) \sin(\omega_i \pi y),$$

with homogeneous Dirichlet boundary conditions and $\omega_i = 2^i$.


For increasing values of n , we obtain the **solutions**:



$n = 1$




$n = 3$




$n = 5$

Multi-Level FBPINNs for a Multi-Frequency Problem – Strong Scaling



Multi-Frequency Problem – What the FBPINN Learns



Cf. Dolean, Heinlein, Mishra, Moseley (2024).


Physics-Informed Randomized Neural Networks (PIRaNNs)

Neural networks

A standard **multilayer perceptron (MLP)** with L hidden layers is a **parametric** model of the form

$$u(x, \theta) = F_{L+1}^A \cdot F_L^{W_L, b_L} \circ \dots \circ F_1^{W_1, b_1}(x),$$

where \mathbf{A} is **linear**, and the i th hidden layer is **nonlinear** $F_i^{W_i, b_i}(x) = \sigma(W_i \cdot x + b_i)$.



In order to optimize the loss function

$$\min_{\theta} \mathcal{L}(\theta),$$


all parameters $\theta = (\mathbf{A}, \mathbf{W}_1, \mathbf{b}_1, \dots, \mathbf{W}_L, \mathbf{b}_L)$ are **trained**.

Randomized neural networks

In **randomized neural networks (RaNNs)** as introduced by **Pao and Takefuji (1992)**,

$$u(x, \mathbf{A}) = F_{L+1}^A \cdot F_L^{W_L, b_L} \circ \dots \circ F_1^{W_1, b_1}(x),$$

the weights in the hidden layers are randomly initialized and **fixed**; only \mathbf{A} is trainable.



The model is **linear** with respect to the trainable parameters \mathbf{A} , and the optimization problem reads

$$\min_{\mathbf{A}} \mathcal{L}(\mathbf{A}).$$

This can **simplify the training process**.


Physics-Informed Randomized Neural Networks (PIRaNNs)

Randomized neural networks

In **randomized neural networks (RaNNs)** as introduced by **Pao and Takefuji (1992)**,

$$u(x, \mathbf{A}) = F_{L+1}^{\mathbf{A}} \cdot F_L^{W_L, b_L} \circ \dots \circ F_1^{W_1, b_1}(x),$$

the weights in the hidden layers are randomly initialized and **fixed**; only \mathbf{A} is trainable.



The model is **linear** with respect to the trainable parameters \mathbf{A} , and the optimization problem reads

$$\min_{\mathbf{A}} \mathcal{L}(\mathbf{A}).$$


This can **simplify the training process**.

Domain decomposition for RaNNs

We employ the FBPINNs approach; cf. **Shang, Heinlein, Mishra, Wang (2025)**. This is closely related to the **random feature method (RFM)** by **Chen, Chi, E, Yang (2022)**. In particular, we solve

$$\mathcal{A}[\sum_{j=1}^J \omega_j u_j(\mathbf{A}_j)](\mathbf{x}_i) = f(\mathbf{x}_i),$$



for $i = 1, \dots, N_{\text{PDE}}$; the boundary conditions are incorporated directly into the u_j .



The hidden weights are randomly initialized, the resulting matrices \mathbf{H} and $\mathbf{H}^\top \mathbf{H}$ are block-sparse.

Preconditioning for Domain Decomposition-Based PIRaNNs

One-level Schwarz preconditioner



Based on an **overlapping domain decomposition**, we define a **one-level Schwarz operator** for $K := H^\top H$

$$M_{\text{OS-1}}^{-1} K = \sum_{i=1}^N R_i^\top K_i^{-1} R_i K,$$

where R_i and R_i^\top are restriction and prolongation operators corresponding to Ω'_i , and $K_i := R_i K R_i^\top$.

Here, the matrix K_i could be singular in which case we use a **pseudo inverse** K_i^+ instead of K_i^{-1} .

We also consider **restricted and scaled additive Schwarz preconditioners**; cf. **Cai, Sarkis (1999)**.

Singular Value Decomposition

As discussed before, on each subdomain Ω_j , the RaNN is


$$\begin{aligned} u_j(x, A_j) &= F_{L+1}^A \cdot F_L^{W_L, b_L} \circ \dots \circ F_1^{W_1, b_1}(x) \\ &= A_j [\Phi_1(x) \quad \dots \quad \Phi_k(x)]^\top, \end{aligned}$$

where k is the width of the last hidden layer and the Φ_l are the randomized basis functions.



Consider a **reduced SVD** $\Phi = U \Sigma V^\top$, where the entries of the matrix are $\Phi_{i,l} = \Phi_l(x_i)$. Then, we consider

$$\hat{u}_j(x, A_j) = A_j \hat{V}^\top [\Phi_1(x) \quad \dots \quad \Phi_k(x)]^\top,$$

where \hat{V}^\top is obtained by omitting the right singular vectors corresponding to small singular values.




Results for the Multi-Frequency Problem ($n=2$)




	$M^{-1} = I$		$M^{-1} = M_{AS}^{-1}$		$M^{-1} = M_{RAS}^{-1}$		$M^{-1} = M_{SAS}^{-1}$	
	iter	e_{L^2}	iter	e_{L^2}	iter	e_{L^2}	iter	e_{L^2}
CG	> 2000	$1.95 \cdot 10^{-2}$	8	$5.03 \cdot 10^{-3}$	—	—	—	—
CGS	> 2000	$2.63 \cdot 10^{-2}$	4	$5.04 \cdot 10^{-3}$	24	$5.03 \cdot 10^{-3}$	6	$5.04 \cdot 10^{-3}$
BICG	> 2000	$1.03 \cdot 10^{-2}$	8	$5.08 \cdot 10^{-3}$	32	$5.05 \cdot 10^{-3}$	11	$5.09 \cdot 10^{-3}$
GMRES	> 2000	$8.68 \cdot 10^{-2}$	13	$5.07 \cdot 10^{-3}$	31	$5.06 \cdot 10^{-3}$	11	$5.08 \cdot 10^{-3}$

4×4 subdomains; DoF = 256; $N = 1600$; $\theta^0 \in \mathcal{U}(-1, 1)$; stop.: $\|M^{-1}r^k\|_{L^2}/\|M^{-1}r^0\|_{L^2} \leq 10^{-5}$


Results for the Multi-Frequency Problem



Multi-level FBPINNs; cf. Dolean, Heinlein, Mishra, Moseley (2024)



DD-PIRaNNs; cf. Shang, Heinlein, Mishra, Wang (2025)



Gauss–Newton Training

Many popular optimizers are based on **gradient descent**:

$$\theta^{(m+1)} = \theta^{(m)} - \alpha \nabla_{\theta} \mathcal{L}(\theta^{(m)}),$$

where $\theta^{(m)}$ are the parameters at iteration m , α is the learning rate, and \mathcal{L} is the loss function; a commonly used example is the **Adam** optimizer ([Kingma \(2017\)](#)).

The **Gauss–Newton method** is based on **Newton's method** but uses the **Gramian** G as an **approximate Jacobian**:

$$\theta^{(m+1)} = \theta^{(m)} - \alpha \mathbf{G}^+(\theta^{(m)}) \nabla_{\theta} \mathcal{L}(\theta^{(m)}),$$

where $G^+(\theta)$ is the pseudoinverse of G , as it may **not** be **invertible** but **only symmetric positive semidefinite**. Therefore, we may regularize G and consider $G + \mu I$ instead, making the matrix **symmetric positive definite**.

Mean squared error (MSE) loss

$$\mathbf{G}(\theta)_{ij} = \sum_{x_i \in \Omega} (\partial_{\theta_i} u_{\theta}(x_i)) (\partial_{\theta_j} u_{\theta}(x_i))$$

Physics-informed loss function


$$\mathbf{G}(\theta)_{ij} = \sum_{x_i \in \Omega} \partial_{\theta_i} n[u]_{\theta}(x_i) \partial_{\theta_j} n[u]_{\theta}(x_i)$$

See, for instance, [Müller and Zeinhofer \(2023\)](#) and [Cai et al. \(arXiv 2024\)](#).

Results for the ODE Problem


Sparsity

- The domain decomposition introduces sparsity in G :
- 8 subdomains
 - coupling blocks due to overlap between subdomains



Cf. Heinlein and Kapoor (arXiv 2025).

Comparison of Adam and Gauss–Newton training




	FBPINN (Adam)	FBPINN (Gauss–Newton)
test MSE	7.8×10^{-3}	8.0×10^{-4}


Spectral analysis and domain decomposition for deep operator networks

Function Versus Operator Learning

Function learning




Operator learning



Deep Operator Networks (DeepONets / DONs)

Neural operators learn operators between function spaces using neural networks. Here, we learn the **solution operator** of a initial-boundary value problem parametrized with p_1, \dots, p_m using **DeepONets** as introduced in [Lu et al. \(2021\)](#).



Single-layer case

The DeepONet architecture is based on the **single-layer case** analyzed in [Chen and Chen \(1995\)](#). In particular, the authors show **universal approximation properties for continuous operators**.

The architecture is based on the following ansatz for presenting the parametrized solution

$$u_{(p_1, \dots, p_m)}(x, t) = \sum_{i=1}^p \underbrace{b_i(p_1, \dots, p_m)}_{\text{branch}} \cdot \underbrace{t_i(x, t)}_{\text{trunk}}$$

Physics-informed DeepONets

DeepONets are compatible with the PINN approach but **physics-informed DeepONets (PI-DeepONets)** are challenging to train.

Other operator learning approaches


- **FNOs:** Li et al. (2021)
- **PCA-Net:** Bhattacharya et al. (2021)
- **Random features:** Nelsen and Stuart (2021)
- **CNOs:** Raonić et al. (2023)

How a DeepONet Maps Between Function Spaces

To illustrate how a DeepONet operates, we consider the **Korteweg–de Vries (KdV) equation**

$$\frac{\partial u}{\partial t} = -u \frac{\partial u}{\partial x} - 0.01 \frac{\partial^3 u}{\partial x^3},$$

which models unidirectional waves in shallow water. Our goal is to train a DeepONet that predicts the wave profile at a future time \hat{t} from the observed height profile $u(\cdot, t_0)$.




Here, the forecast time \hat{t} is fixed to keep the learning task simple. A **more general neural operator** can take the target time as an additional input to the trunk network.

How a DeepONet Maps Between Function Spaces

To illustrate how a DeepONet operates, we consider the **Korteweg–de Vries (KdV)** equation

$$\frac{\partial u}{\partial t} = -u \frac{\partial u}{\partial x} - 0.01 \frac{\partial^3 u}{\partial x^3},$$

which models unidirectional waves in shallow water. Our goal is to train a DeepONet that predicts the wave profile at a future time \hat{t} from the observed height profile $u(\cdot, t_0)$.




Here, the forecast time \hat{t} is fixed to keep the learning task simple. A **more general neural operator** can take the target time as an additional input to the trunk network.

DeepONet Trunk Basis – Examples

Let us consider some examples of the left singular vectors for three differential equations:

- A) advection-diffusion equation
- B) KdV equation
- C) Burgers equation




The learned trunk bases have been investigated in more detail in [Williams et al. \(2024\)](#).

DeepONet – Error Decomposition Results for the KdV Equation

Results for the **KdV equation** with $t_0 = 0.0$ and $\hat{t} = 0.2$.

900 training and 100 test configurations.



total error
 δ

trunk error
 δ_T

branch error
 δ_B

SVD truncation error
 δ_{SVD}

DeepONet – Branch Error

Using the left singular vectors, the **branch error** becomes

$$\mathcal{E}_B = \sum_{i=1}^m \sigma_i^2 \underbrace{\|b_i - v_i\|_2^2}_{=:L_i}.$$


We call

$$\sigma_i^2 L_i$$

the **(weighted) mode loss** because it equals the loss contribution of the i th mode. Accordingly, L_i is the **unweighted mode loss**.

This choice of **left singular vectors as the trunk basis** is often denoted **POD-DeepONet** in **Lu et al. (2022)**.

Unweighted mode loss



DeepONet – Branch Error

Using the left singular vectors, the **branch error** becomes

$$\mathcal{E}_B = \sum_{i=1}^m \sigma_i^2 \underbrace{\|b_i - v_i\|_2^2}_{=:L_i}.$$


We call

$$\sigma_i^2 L_i$$

the **(weighted) mode loss** because it equals the loss contribution of the i th mode. Accordingly, L_i is the **unweighted mode loss**.

This choice of **left singular vectors as the trunk basis** is often denoted **POD-DeepONet** in **Lu et al. (2022)**.

Unweighted mode loss



DeepONet – Branch Error

Using the left singular vectors, the **branch error** becomes

$$\mathcal{E}_B = \sum_{i=1}^m \sigma_i^2 \underbrace{\|b_i - v_i\|_2^2}_{=:L_i}.$$


We call

$$\sigma_i^2 L_i$$

the **(weighted) mode loss** because it equals the loss contribution of the i th mode. Accordingly, L_i is the **unweighted mode loss**.

This choice of **left singular vectors as the trunk basis** is often denoted **POD-DeepONet** in **Lu et al. (2022)**.

Unweighted mode loss



DeepONet – Branch Error

Using the left singular vectors, the **branch error** becomes

$$\mathcal{E}_B = \sum_{i=1}^m \sigma_i^2 \underbrace{\|b_i - v_i\|_2^2}_{=:L_i}.$$


We call

$$\sigma_i^2 L_i$$

the **(weighted) mode loss** because it equals the loss contribution of the i th mode. Accordingly, L_i is the **unweighted mode loss**.

This choice of **left singular vectors as the trunk basis** is often denoted **POD-DeepONet** in **Lu et al. (2022)**.

Unweighted mode loss



DeepONet – Branch Error

Using the left singular vectors, the **branch error** becomes

$$\mathcal{E}_B = \sum_{i=1}^m \sigma_i^2 \underbrace{\|b_i - v_i\|_2^2}_{=:L_i}.$$


We call

$$\sigma_i^2 L_i$$

the **(weighted) mode loss** because it equals the loss contribution of the i th mode. Accordingly, L_i is the **unweighted mode loss**.

This choice of **left singular vectors as the trunk basis** is often denoted **POD-DeepONet** in **Lu et al. (2022)**.

Unweighted mode loss



→ The coefficients of modes with **large singular values** are **learned best**. Errors for modes with **small singular values** **remain high**.

DeepONet – Branch Error

Using the left singular vectors, the **branch error** becomes

$$\mathcal{E}_B = \sum_{i=1}^m \sigma_i^2 \underbrace{\|b_i - v_i\|_2^2}_{=:L_i}.$$


We call

$$\sigma_i^2 L_i$$

the **(weighted) mode loss** because it equals the loss contribution of the i th mode. Accordingly, L_i is the **unweighted mode loss**.

This choice of **left singular vectors as the trunk basis** is often denoted **POD-DeepONet** in **Lu et al. (2022)**.

(Weighted) mode loss



DeepONet – Branch Error

Using the left singular vectors, the **branch error** becomes

$$\mathcal{E}_B = \sum_{i=1}^m \sigma_i^2 \underbrace{\|b_i - v_i\|_2^2}_{=:L_i}.$$

We call


$$\sigma_i^2 L_i$$

the **(weighted) mode loss** because it equals the loss contribution of the i th mode. Accordingly, L_i is the **unweighted mode loss**.

This choice of **left singular vectors as the trunk basis** is often denoted **POD-DeepONet** in **Lu et al. (2022)**.

→ Analyzing the actual error contributions, the **modes with medium singular values contribute most**.

(Weighted) mode loss



DeepONet – Branch Error

Using the left singular vectors, the **branch error** becomes

$$\mathcal{E}_B = \sum_{i=1}^m \sigma_i^2 \underbrace{\|b_i - v_i\|_2^2}_{=:L_i}.$$


We call

$$\sigma_i^2 L_i$$

the **(weighted) mode loss** because it equals the loss contribution of the i th mode. Accordingly, L_i is the **unweighted mode loss**.


This choice of **left singular vectors as the trunk basis** is often denoted **POD-DeepONet** in **Lu et al. (2022)**.

(Weighted) mode loss



How to improve the performance on medium-sized singular value modes?

Finite Basis DeepONets (FBDONs)



Howard, Heinlein, Stinis (in prep.)

Variants:

Shared-trunk FBDONs (ST-FBDONs)

The trunk net learns spatio-temporal basis functions. In ST-FBDONs, we use the **same trunk network for all subdomains**.

Stacking FBDONs

Combination of the **stacking multifidelity approach** with FBDONs.

Heinlein, Howard, Beecroft, Stinis (2025)


FBDONs – Wave Equation

Wave equation

$$\frac{d^2s}{dt^2} = 2 \frac{d^2s}{dx^2}, \quad (x, t) \in [0, 1]^2$$

$$s_t(x, 0) = 0, x \in [0, 1], \quad s(0, t) = s(1, t) = 0,$$

Solution: $s(x, t) = \sum_{n=1}^5 b_n \sin(n\pi x) \cos(n\pi\sqrt{2}t)$



Parametrization

Initial conditions for s parametrized by $b = (b_1, \dots, b_5)$ (normally distributed):

$$s(x, 0) = \sum_{n=1}^5 b_n \sin(n\pi x) \quad x \in [0, 1]$$

Training on 1 000 random configurations.

Mean rel. ℓ_2 error on 100 config.


	Mean rel. ℓ_2 error on 100 config.
DeepONet	0.30 ± 0.11
ML-ST-FBDON ([1, 4, 8, 16] subd.)	0.05 ± 0.03
ML-FBDON ([1, 4, 8, 16] subd.)	0.08 ± 0.04

→ Sharing the trunk network does not only save in the number of parameters but even yields **better performance**

Cf. [Howard, Heinlein, Stinis \(in prep.\)](#)

Domain decomposition-based image segmentation for high-resolution image segmentation on multiple GPUs

Memory Requirements for CNN Training




- As an example for a **convolutional neural network (CNN)**, we employ the **U-Net architecture** introduced in **Ronneberger, Fischer, and Brox (2015)**.
- The U-Net yields **state-of-the-art accuracy in semantic image segmentation** and other **image-to-image tasks**.

Below: memory consumption for training on a single 1024×1024 image.


name	size	# channels		mem. feature maps		mem. weights	
		input	output	# of values	MB	# of values	MB
input block	1 024	3	64	268 M	1 024.0	38 848	0.148
encoder block 1	512	64	128	167 M	704.0	221 696	0.846
encoder block 2	256	128	256	84 M	352.0	885 760	3.379
encoder block 3	128	256	512	42 M	176.0	3 540 992	13.508
encoder block 4	64	512	1 024	21 M	88.0	14 159 872	54.016
decoder block 1	64	1,024	512	50 M	192.0	9 177 088	35.008
decoder block 2	128	512	256	101 M	384.0	2 294 784	8.754
decoder block 3	256	256	128	201 M	768.0	573 952	2.189
decoder block 4	512	128	64	402 M	1 536.0	143 616	0.548
output block	1 024	64	3	3.1 M	12.0	195	0.001

Decomposing the U-Net




Cf. Verburg, Heinlein, Cyr (2025).


Decomposing the U-Net



Decomposing the U-Net




Decomposing the U-Net



- Distribution of feature maps results in **significant reduction of memory usage on a single GPU**
- Moderate **additional memory usage** due to the **communication network**

Results – Synthetic Data Set

Task: Connect two dots via a line segment



Result: Communication




True mask



Pred. (no comm.)




Pred. (comm.)



DeepGlobe 2018 Satellite Image Data Set (Demir et al. (2018))


class	pixel count	proportion
urban	642.4M	9.35 %
agriculture	3898.0M	56.76 %
rangeland	701.1M	10.21 %
forest	944.4M	13.75 %
water	256.9M	3.74 %
barren	421.8M	6.14 %
unknown	3.0M	0.04 %



Avoiding overfitting


The data set includes **only 803 images**. To **avoid overfitting**, we

- apply **batch normalization**, use **random dropout** layers and **data augmentation**, and
- initialize the encoder using the **ResNet-18** (**He, Zhang, Ren, and Sun (2016)**)



DeepGlobe 2018 Satellite Image Data Set (Demir et al. (2018))


class	pixel count	proportion
urban	642.4M	9.35 %
agriculture	3898.0M	56.76 %
rangeland	701.1M	10.21 %
forest	944.4M	13.75 %
water	256.9M	3.74 %
barren	421.8M	6.14 %
unknown	3.0M	0.04 %



Avoiding overfitting

The data set includes **only 803 images**. To **avoid overfitting**, we

- apply **batch normalization**, use **random dropout** layers and **data augmentation**, and
- initialize the encoder using the **ResNet-18** (He, Zhang, Ren, and Sun (2016))



CWI Research Semester Programme:

Bridging Numerical Analysis and Scientific Machine Learning: Advances and Applications

Co-organizers: Victorita Dolean (TU/e), Alexander Heinlein (TU Delft), Benjamin Sanderse (CWI), Jemima Tabbeart (TU/e), Tristan van Leeuwen (CWI)

- **Autumn School** (October 27–31, 2025):

- Chris Budd (University of Bath)
- Ben Moseley (Imperial College London)
- Gabriele Steidl (Technische Universität Berlin)
- Andrew Stuart (California Institute of Technology)
- Andrea Walther (Humboldt-Universität zu Berlin)
- Ricardo Baptista (University of Toronto)

- **Workshop** (December 1–3, 2025):

- Plenary talks (academia & industry) and panel discussion
- **Poster session with prize sponsored by Math4NL**
- Plenary speakers:
 - Benjamin Peherstorfer (NYU)
 - Elena Celledoni (NTNU)
 - Jakob Sauer Jørgensen (DTU)
 - Marcelo Pereyra (Heriot-Watt University)
 - Nicolas Boullé (ICL)



Centrum Wiskunde & Informatica



Join us for inspiring talks, hands-on sessions, and industry collaboration!

Summary


Domain Decomposition-Based Neural Network Architectures

- Localization via domain decomposition lets subnetworks learn local features, mitigates the spectral bias in PDE surrogates, and captures sharp variations more accurately.
- Partitioning the domain yields parallelism, promotes sparsity, and keeps the architecture computationally efficient.

Domain Decomposition Preconditioning

- The DD architecture behaves like a discretization, so we precondition the least-squares system with well-known algorithms.
- One-level DD preconditioners reduces large eigenvalues, while SVD-based reduction removes near-linear dependencies.

Thank you for your attention!



Topical Activity
Group
Scientific Machine
Learning

