



Neural Network-Based Models for Physical Systems


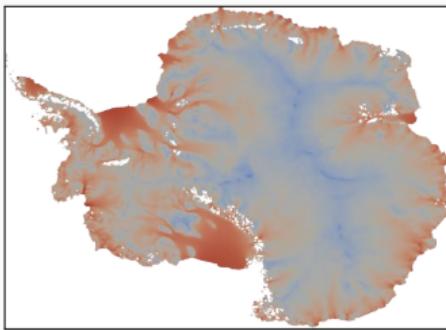
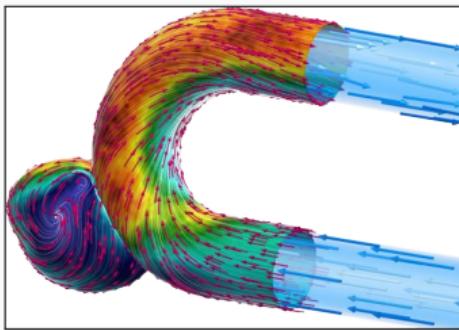
Analysis, Domain Decomposition, and Preconditioning

Alexander Heinlein¹

International Conference on Latest Advances in Computational and Applied Mathematics 2025 (LACAM-25), IISER Thiruvananthapuram, Kerala, India, December 8-11, 2025

¹Delft University of Technology

Scientific Computing and Machine Learning



Numerical methods

Based on physical models

- + Robust and generalizable
- Require availability of mathematical models

Machine learning models

Driven by data

- + Do not require mathematical models
- Sensitive to data, limited extrapolation capabilities

Scientific machine learning

Combining the strengths and compensating the weaknesses of the individual approaches:

numerical methods	improve	machine learning techniques
machine learning techniques	assist	numerical methods

Outline

1 Physics-informed neural networks – Adaptive sampling and localization via domain decomposition

Based on joint work with

Damien Beecroft

(University of Washington)

Victorita Dolean

(Eindhoven University of Technology)

Bianca Giovanardi, Coen Visser

(Delft University of Technology)

Amanda A. Howard and Panos Stinis

(Pacific Northwest National Laboratory)

Siddhartha Mishra

(ETH Zürich)

Ben Moseley

(Imperial College London)

2 Deep operator networks – Error analysis and localization via domain decomposition

Based on joint work with

Damien Beecroft

(University of Washington)

Amanda A. Howard and Panos Stinis

(Pacific Northwest National Laboratory)

Johannes Taraz

(Delft University of Technology)

Physics-informed neural networks – Adaptive sampling and localization via domain decompostion

Physics-Informed Neural Networks (PINNs)

In the **physics-informed neural network (PINN)** approach introduced by **Raissi et al. (2019)**, a **neural network** is employed to **discretize a partial differential equation**

$$\mathcal{N}[u] = f, \quad \text{in } \Omega.$$

PINNs use a **hybrid loss function**:



$$\mathcal{L}(\theta) = \omega_{\text{data}} \mathcal{L}_{\text{data}}(\theta) + \omega_{\text{PDE}} \mathcal{L}_{\text{PDE}}(\theta),$$

where ω_{data} and ω_{PDE} are **weights** and


$$\mathcal{L}_{\text{data}}(\theta) = \frac{1}{N_{\text{data}}} \sum_{i=1}^{N_{\text{data}}} (u(\hat{x}_i, \theta) - u_i)^2,$$

$$\mathcal{L}_{\text{PDE}}(\theta) = \frac{1}{N_{\text{PDE}}} \sum_{i=1}^{N_{\text{PDE}}} (\mathcal{N}[u](x_i, \theta) - f(x_i))^2.$$

See also Dissanayake and Phan-Thien (1994); Lagaris et al. (1998).



Hybrid loss



Advantages

- "Meshfree"
- Small data
- Generalization properties
- High-dimensional problems
- Inverse and parameterized problems

Drawbacks

- Training cost and robustness
- Convergence not well-understood
- Difficulties with scalability and multi-scale problems

- Known solution values can be included in $\mathcal{L}_{\text{data}}$
- Initial and boundary conditions are also included in $\mathcal{L}_{\text{data}}$

Estimate of the total error (Mishra and Molinaro (2022))

The total error (or generalization error) satisfies

$$\mathcal{E}_G \leq C_{\text{PDE}} \mathcal{E}_T + C_{\text{PDE}} C_{\text{quad}}^{1/p} N^{-\alpha/p}$$


- $\mathcal{E}_G = \mathcal{E}_G(\mathbf{X}, \theta) := \|\mathbf{u} - \mathbf{u}^*\|_V$ **general. error** (V Sobolev space, \mathbf{X} training data set)
 - \mathcal{E}_T **training error** (l^p loss of the residual of the PDE)
 - N **number of the training points** and α **convergence rate of the quadrature**
 - C_{PDE} and C_{quad} **constants** depending on the **PDE**, quadrature, and neural network
-
- **Training error:** optimization is **challenging** due to the **highly nonconvex** loss landscape of PINNs.
 - **Sampling error:** unlike classical numerical methods, standard quadrature rules **poorly approximate** neural-network trial functions.

PACMANN – Point Adaptive Collocation Method for Artificial Neural Networks

In **Visser, Heinlein, and Giovanardi (subm. 2025; arXiv:2411.19632)**,
the collocation points are updated by solving the **min-max problem**

$$\min_{\theta} \left[\omega_{\text{data}} \mathcal{L}_{\text{data}}(\theta) + \max_{\mathbf{X} \subset \Omega} \omega_{\text{PDE}} \mathcal{L}_{\text{PDE}}(\mathbf{X}, \theta) \right].$$

This idea was already mentioned in **Wang et al. (2022)**. Different from other residual-based adaptive sampling methods, **existing collocation points are moved** using gradient-based optimizers such as **gradient ascent, RMSprop (Hinton (2018))**, or **Adam (Kingma, Ba (2017))**.



Algorithm 1: PACMANN with iteration counts P and T and stepsize s

Sample a set \mathbf{X} of N_{PDE} collocation points using a uniform sampling method;

while stopping criterion not reached **do**

 Train the PINN for P iterations;

for $k = 1, \dots, T$ **do**

 Compute squared residual $\mathcal{R}(x_i) = (\mathcal{N}[u](x_i, \theta) - f(x_i))^2$ for all $x_i \in \mathbf{X}$;

 Compute gradient $\nabla_{\mathbf{x}} \mathcal{R}(x_i)$ for all $x_i \in \mathbf{X}$;

 Move the points in \mathbf{X} according to the chosen optimization algorithm and stepsize s ;

end

 Resample points in \mathbf{X} that moved outside Ω based on a uniform probability distribution;


end

PACMANN – Point Adaptive Collocation Method for Artificial Neural Networks

In [Visser, Heinlein, and Giovanardi \(subm. 2025; arXiv:2411.19632\)](#),
the collocation points are updated by solving the **min-max problem**


$$\min_{\theta} \left[\omega_{\text{data}} \mathcal{L}_{\text{data}}(\theta) + \max_{\mathbf{X} \subset \Omega} \omega_{\text{PDE}} \mathcal{L}_{\text{PDE}}(\mathbf{X}, \theta) \right].$$

This idea was already mentioned in [Wang et al. \(2022\)](#). Different from other residual-based adaptive sampling methods, **existing collocation points are moved** using gradient-based optimizers such as [gradient ascent](#), [RMSprop](#) ([Hinton \(2018\)](#)), or [Adam](#) ([Kingma, Ba \(2017\)](#)).



Burgers' equation – Comparison against different methods

sampling method <i>2500 coll. points</i>	L_2 relative error		mean runtime [s]
	mean	1 SD	
uniform grid	25.9%	14.2%	425
Hammersley grid	0.61%	0.53%	443
random resampling	0.40%	0.35%	423
RAR	0.11%	0.05%	450
RAD	0.16%	0.10%	463
RAR-D	0.24%	0.21%	503
PACMANN–Adam	0.07%	0.05%	461





PACMANN – Point Adaptive Collocation Method for Artificial Neural Networks

In [Visser, Heinlein, and Giovanardi \(subm. 2025; arXiv:2411.19632\)](#),
the collocation points are updated by solving the **min-max problem**


$$\min_{\theta} \left[\omega_{\text{data}} \mathcal{L}_{\text{data}}(\theta) + \max_{X \subset \Omega} \omega_{\text{PDE}} \mathcal{L}_{\text{PDE}}(X, \theta) \right].$$

This idea was already mentioned in [Wang et al. \(2022\)](#). Different from other residual-based adaptive sampling methods, **existing collocation points are moved** using gradient-based optimizers such as [gradient ascent](#), [RMSprop](#) ([Hinton \(2018\)](#)), or [Adam](#) ([Kingma, Ba \(2017\)](#)).



Spectral Bias of Neural Networks

Rahaman et al. (2019) observed the **spectral bias of neural networks**: during training, they learn low-frequency functions faster than high-frequency functions; see also Cao et al. (2021), Wang, et al. (2022), Hong et al. (arXiv 2022), Xu et al. (2024), ...



This can be understood by interpreting the training dynamics of neural networks as **gradient flow**

$$\theta_{k+1} = \theta_k - \eta_k \nabla_\theta \mathcal{C}(n_{\theta_k}) \quad \rightarrow \quad \frac{dx}{dt}(t) = -\nabla_x \mathcal{C}(x(t))$$


For **infinitely wide neural networks**, we obtain a **constant neural tangent kernel (NTK)** (Jacot et al. (2018)).

For mean squared error (MSE) loss \mathcal{C} , we obtain the **discretized K** = $\left(\left\langle \frac{dn_{\theta(t)}}{d\theta}(x_i), \frac{dn_{\theta(t)}}{d\theta}(x_j) \right\rangle \right)_{ij}$:

$$\frac{dn_{\theta(t)}}{dt} = -\frac{2}{n} \cdot K(t) (n_{\theta(t)}(\mathbf{X}) - \mathbf{Y}) \quad \Rightarrow \quad \mathbf{U}^\top (n_{\theta(t)}(\mathbf{X}) - \mathbf{Y}) \approx e^{-t \frac{2}{n} \Lambda} \mathbf{U}^\top \mathbf{Y}$$

Spectral Bias of Neural Networks

Rahaman et al. (2019) observed the **spectral bias of neural networks**: during training, they learn **low-frequency functions faster than high-frequency functions**; see also Cao et al. (2021), Wang, et al. (2022), Hong et al. (arXiv 2022), Xu et al. (2024), ...



Scaling of PINNs for a Simple ODE Problem

Solve


$$\begin{aligned} u' &= \cos(\omega x), \\ u(0) &= 0, \end{aligned}$$

for different values of ω
using PINNs with
varying network
capacities.

Scaling issues

- Large computational domains
- Small frequencies

Cf. Moseley, Markham, and
Nissen-Meyer (2023)



Scaling of PINNs for a Simple ODE Problem

Solve


$$\begin{aligned} u' &= \cos(\omega x), \\ u(0) &= 0, \end{aligned}$$

for different values of ω
using PINNs with
varying network
capacities.

Scaling issues

- Large computational domains
- Small frequencies

Cf. Moseley, Markham, and Nissen-Meyer (2023)



Domain Decomposition Methods and Machine Learning – Literature

A non-exhaustive literature overview:

- ML for adaptive BDDC, FETI–DP, AGDSW: H., Klawonn, Langer, Weber (2019, 2020, 2021, 2021, 2021, 2022); Klawonn, Langer, Weber (2024, 2025)
- cPINNs, XPINNs: Jagtap, Kharazmi, Karniadakis (2020); Jagtap, Karniadakis (2020)
- Classical Schwarz iteration for PINNs or DeepRitz:: Li, Tang, Wu, and Liao (2019); Li, Xiang, Xu (2020); Mercier, Gratton, Boudier (arXiv 2021); Dolean, H., Mercier, Gratton (acc. 2025); Li, Wang, Cui, Xiang, Xu (2023); Sun, Xu, Yi (arXiv 2023, 2024); Kim, Yang (2023, 2024, 2024)
- FBPINNs, FBKANs: Moseley, Markham, Nissen-Meyer (2023); Dolean, H., Mishra, Moseley (2024, 2024); H., Howard, Beecroft, Stinis (2025); Howard, Jacob, Murphy, H., Stinis (arXiv 2024)
- DD for randomized NNs: Dong, Li (2021); Dang, Wang (2024); Sun, Dong, Wang (2024); Sun, Wang (2024); Chen, Chi, E, Yang (2022); Shang, H., Mishra, Wang (2025); Anderson, Dolean, Moseley, Pestana, (arXiv 2024); van Beek, Dolean, Moseley (arxiv 2025)
- DD for Neural Operators and Surrogate Models: H., Howard, Beecroft, Stinis (2025); Ramezankhani, Parekh, Deodhar, Birru (arXiv 2024); Wu, Kovachki, Liu (arXiv 2025); Pelzer, Verburg, H., Schulte (arXiv 2025); Klaes, Klawonn, Kubicki, Langer, Nakajima, Shimokawabe, Weber (arXiv 2025); Howard, H., Stinis (in prep.)
- DD for CNNs: Gu, Zhang, Liu, Cai (2022); Lee, Park, Lee (2022); Klawonn, Langer, Weber (2024); Verburg, H., Cyr (2025)

An overview of the state-of-the-art in 2024:



A. Klawonn, M. Langer, J. Weber

Machine learning, domain decomposition methods – a survey

Computational Science and Engineering. 2024

Finite Basis Physics-Informed Neural Networks (FBPINNs)


FBPINNs ([Moseley, Markham, Nissen-Meyer \(2023\)](#))

FBPINNs employ the **network architecture**


$$u(\theta_1, \dots, \theta_J) = \sum_{j=1}^J \omega_j u_j(\theta_j)$$

and the **loss function**


$$\mathcal{L} = \frac{1}{N} \sum_{i=1}^N \left(n \left[\sum_{x_i \in \Omega_j} \omega_j u_j(x_i, \theta_j) - f(x_i) \right] \right)^2$$



1D single-frequency problem



[Moseley, Markham, Nissen-Meyer \(2023\)](#)



Finite Basis Physics-Informed Neural Networks (FBPINNs)


FBPINNs ([Moseley, Markham, Nissen-Meyer \(2023\)](#))

FBPINNs employ the **network architecture**


$$u(\theta_1, \dots, \theta_J) = \sum_{j=1}^J \omega_j u_j(\theta_j)$$

and the **loss function**

$$\mathcal{L} = \frac{1}{N} \sum_{i=1}^N \left(n \left[\sum_{x_i \in \Omega_j} \omega_j u_j(x_i, \theta_j) - f(x_i) \right] \right)^2$$




1D single-frequency problem



Multi-Level FBPINNs

Multi-level FBPINNs (ML-FBPINNs)

ML-FBPINNs ([Dolean, Heinlein, Mishra, Moseley \(2024\)](#)) are based on a **hierarchy** of domain decompositions:




This yields the **network architecture**:


$$u(\theta_1^{(1)}, \dots, \theta_{J^{(L)}}^{(L)}) = \sum_{l=1}^L \sum_{j=1}^{N^{(l)}} \omega_j^{(l)} u_j^{(l)}(\theta_j^{(l)})$$

Multiscale problems ...


... appear in most areas of modern science and engineering:



Dual-phase steel;
fig. courtesy of
[J. Schröder](#).



Groundwater flow;
cf. [Christie & Blunt \(2001\)](#) (SPE10).



Arterial walls;
cf. [O'Connell et al. \(2008\)](#).

Multi-frequency problem

Consider the **multi-frequency Laplace problem**

$$-\Delta u = 2 \sum_{i=1}^n (\omega_i \pi)^2 \sin(\omega_i \pi x) \sin(\omega_i \pi y),$$

with homogeneous Dirichlet boundary conditions and $\omega_i = 2^i$.

For increasing values of n , we obtain the **solutions**:



$n = 1$



$n = 3$




$n = 5$

Multi-Level FBPINNs

Multi-level FBPINNs (ML-FBPINNs)

ML-FBPINNs ([Dolean, Heinlein, Mishra, Moseley \(2024\)](#)) are based on a **hierarchy** of domain decompositions:




This yields the **network architecture**:


$$u(\theta_1^{(1)}, \dots, \theta_{J^{(L)}}^{(L)}) = \sum_{l=1}^L \sum_{i=1}^{N^{(l)}} \omega_j^{(l)} u_j^{(l)}(\theta_j^{(l)})$$

Multiscale problems ...


... appear in most areas of modern science and engineering:



Dual-phase steel;
fig. courtesy of
[J. Schröder](#).



Groundwater flow;
cf. [Christie & Blunt \(2001\)](#) (SPE10).



Arterial walls;
cf. [O'Connell et al. \(2008\)](#).




Multi-frequency problem

Consider the **multi-frequency Laplace problem**


$$-\Delta u = 2 \sum_{i=1}^n (\omega_i \pi)^2 \sin(\omega_i \pi x) \sin(\omega_i \pi y),$$

with homogeneous Dirichlet boundary conditions and $\omega_i = 2^i$.


For increasing values of n , we obtain the **solutions**:



Multi-Level FBPINNs for a Multi-Frequency Problem – Strong Scaling



Multi-Frequency Problem – What the FBPINN Learns




Cf. Dolean, Heinlein, Mishra, Moseley (2024).

Multifidelity Stacking FBPINNs – Allen–Cahn Equation

Finally, we consider the **Allen–Cahn equation**, describing phase separation in multi-component alloy systems:

$$\begin{aligned}\varphi_t - 0.0001\varphi_{xx} + 5\varphi^3 - 5\varphi &= 0, & t \in (0, 1], x \in [-1, 1], \\ \varphi(x, 0) &= x^2 \cos(\pi x), & x \in [-1, 1], \\ \varphi(x, t) &= \varphi(-x, t), & t \in [0, 1], x = -1, x = 1, \\ \varphi_x(x, t) &= \varphi_x(-x, t), & t \in [0, 1], x = -1, x = 1.\end{aligned}$$



PINN gets stuck at fixed point of the dynamical system; cf. Rohrhofer et al. (2023).


Domain Decomposition for Physics-Informed Randomized Neural Networks

Neural networks

A standard **multilayer perceptron (MLP)** with L hidden layers is a **parametric** model of the form

$$u(x, \theta) = F_{L+1}^A \cdot F_L^{W_L, b_L} \circ \dots \circ F_1^{W_1, b_1}(x),$$

where A is **linear**, and the i th hidden layer is **nonlinear** $F_i^{W_i, b_i}(x) = \sigma(W_i \cdot x + b_i)$.



In order to optimize the loss function

$$\min_{\theta} \mathcal{L}(\theta),$$


all parameters $\theta = (A, W_1, b_1, \dots, W_L, b_L)$ are **trained**.

Randomized neural networks

In **randomized neural networks (RaNNs)** as introduced by **Pao and Takefuji (1992)**,

$$u(x, A) = F_{L+1}^A \cdot F_L^{W_L, b_L} \circ \dots \circ F_1^{W_1, b_1}(x),$$

the weights in the hidden layers are randomly initialized and **fixed**; only A is trainable.



The model is **linear** with respect to the trainable parameters A , and the optimization problem reads

$$\min_A \mathcal{L}(A).$$

This can **simplify the training process**.


Domain Decomposition for Physics-Informed Randomized Neural Networks

Domain decomposition for RaNNs

We employ the FBPINNs approach; cf. **Shang, Heinlein, Mishra, Wang (2025)**. This is closely related to the **random feature method (RFM)** by **Chen, Chi, E, Yang (2022)**. In particular, we solve

$$\mathcal{A}[\sum_{j=1}^J \omega_j u_j(\mathbf{A}_j)](\mathbf{x}_i) = f(\mathbf{x}_i),$$

for $i = 1, \dots, N_{\text{PDE}}$; the boundary conditions are incorporated directly into the u_j .




The hidden weights are randomly initialized, the resulting matrices \mathbf{H} and $\mathbf{H}^\top \mathbf{H}$ are block-sparse.

Randomized neural networks

In **randomized neural networks (RaNNs)** as introduced by **Pao and Takefuji (1992)**,

$$u(\mathbf{x}, \mathbf{A}) = F_{L+1}^{\mathbf{A}} \cdot F_L^{W_L, b_L} \circ \dots \circ F_1^{W_1, b_1}(\mathbf{x}),$$

the weights in the hidden layers are randomly initialized and **fixed**; only \mathbf{A} is trainable.



The model is **linear** with respect to the trainable parameters \mathbf{A} , and the optimization problem reads

$$\min_{\mathbf{A}} \mathcal{L}(\mathbf{A}).$$

This can **simplify the training process**.


Domain Decomposition for Physics-Informed Randomized Neural Networks

Domain decomposition for RaNNs

We employ the FBPINNs approach; cf. **Shang, Heinlein, Mishra, Wang (2025)**. This is closely related to the **random feature method (RFM)** by **Chen, Chi, E, Yang (2022)**. In particular, we solve


$$\mathcal{A}[\sum_{j=1}^J \omega_j u_j(\mathbf{A}_j)](\mathbf{x}_i) = f(\mathbf{x}_i),$$

for $i = 1, \dots, N_{\text{PDE}}$; the boundary conditions are incorporated directly into the u_j .




The hidden weights are randomly initialized, the resulting matrices \mathbf{H} and $\mathbf{H}^\top \mathbf{H}$ are block-sparse.

One-level Schwarz preconditioner



Solution of local problem



$$\mathbf{M}_{\text{OS-1}}^{-1} \mathbf{K} = \sum_{i=1}^N \mathbf{R}_i^\top \mathbf{K}_i^{-1} \mathbf{R}_i \mathbf{K},$$

where \mathbf{R}_i restriction operator to Ω'_i , $\mathbf{K}_i := \mathbf{R}_i \mathbf{K} \mathbf{R}_i^\top$.

Singular Value Decomposition (SVD)

On each subdomain Ω_j , the RaNN reads



$$u_j(\mathbf{x}, \mathbf{A}_j) = \mathbf{A}_j [\Phi_1(\mathbf{x}) \quad \dots \quad \Phi_k(\mathbf{x})]^\top,$$

where Φ_1, \dots, Φ_k are the randomized basis functions in the **last hidden layer**. Replace it by

$$\hat{u}_j(\mathbf{x}, \mathbf{A}_j) = \mathbf{A}_j \hat{\mathbf{V}}^\top [\Phi_1(\mathbf{x}) \quad \dots \quad \Phi_k(\mathbf{x})]^\top,$$

where $\hat{\mathbf{V}}^\top$ truncated right singular vectors of Φ .


Results for the Multi-Frequency Problem ($n=2$)




	$M^{-1} = I$		$M^{-1} = M_{AS}^{-1}$		$M^{-1} = M_{RAS}^{-1}$		$M^{-1} = M_{SAS}^{-1}$	
	iter	e_{L^2}	iter	e_{L^2}	iter	e_{L^2}	iter	e_{L^2}
CG	> 2000	$1.95 \cdot 10^{-2}$	8	$5.03 \cdot 10^{-3}$	—	—	—	—
CGS	> 2000	$2.63 \cdot 10^{-2}$	4	$5.04 \cdot 10^{-3}$	24	$5.03 \cdot 10^{-3}$	6	$5.04 \cdot 10^{-3}$
BICG	> 2000	$1.03 \cdot 10^{-2}$	8	$5.08 \cdot 10^{-3}$	32	$5.05 \cdot 10^{-3}$	11	$5.09 \cdot 10^{-3}$
GMRES	> 2000	$8.68 \cdot 10^{-2}$	13	$5.07 \cdot 10^{-3}$	31	$5.06 \cdot 10^{-3}$	11	$5.08 \cdot 10^{-3}$

4×4 subdomains; DoF = 256; $N = 1600$; $\theta^0 \in \mathcal{U}(-1, 1)$; stop.: $\|M^{-1}r^k\|_{L^2}/\|M^{-1}r^0\|_{L^2} \leq 10^{-5}$


Results for the Multi-Frequency Problem



Multi-level FBPINNs; cf. Dolean, Heinlein, Mishra, Moseley (2024)




DD-PIRaNNs; cf. Shang, Heinlein, Mishra, Wang (2025)



Deep operator networks – Error analysis and localization via domain decompostion

Deep Operator Networks (DeepONets / DONs)

Neural operators learn operators between function spaces using neural networks. Here, we learn the **solution operator** of a initial-boundary value problem parametrized with p_1, \dots, p_m using **DeepONets** as introduced in [Lu et al. \(2021\)](#).



Single-layer case

The DeepONet architecture is based on the **single-layer case** analyzed in [Chen and Chen \(1995\)](#). In particular, the authors show **universal approximation properties for continuous operators**.

The architecture is based on the following ansatz for presenting the parametrized solution

$$u_{(p_1, \dots, p_m)}(x, t) = \sum_{i=1}^p \underbrace{b_i(p_1, \dots, p_m)}_{\text{branch}} \cdot \underbrace{t_i(x, t)}_{\text{trunk}}$$

Physics-informed DeepONets

DeepONets are compatible with the PINN approach but **physics-informed DeepONets (PI-DeepONets)** are challenging to train.

Other operator learning approaches


- **FNOs:** Li et al. (2021)
- **PCA-Net:** Bhattacharya et al. (2021)
- **Random features:** Nelsen and Stuart (2021)
- **CNOs:** Raonić et al. (2023)

How a DeepONet Maps Between Function Spaces

To illustrate how a DeepONet operates, we consider the **Korteweg–de Vries (KdV) equation**

$$\frac{\partial u}{\partial t} = -u \frac{\partial u}{\partial x} - 0.01 \frac{\partial^3 u}{\partial x^3},$$

which models unidirectional waves in shallow water. Our goal is to train a DeepONet that predicts the wave profile at a future time \hat{t} from the observed height profile $u(\cdot, t_0)$.




Here, the forecast time \hat{t} is fixed to keep the learning task simple. A **more general neural operator** can take the target time as an additional input to the trunk network.

How a DeepONet Maps Between Function Spaces

To illustrate how a DeepONet operates, we consider the **Korteweg–de Vries (KdV) equation**

$$\frac{\partial u}{\partial t} = -u \frac{\partial u}{\partial x} - 0.01 \frac{\partial^3 u}{\partial x^3},$$

which models unidirectional waves in shallow water. Our goal is to train a DeepONet that predicts the wave profile at a future time \hat{t} from the observed height profile $u(\cdot, t_0)$.



Here, the forecast time \hat{t} is fixed to keep the learning task simple. A **more general neural operator** can take the target time as an additional input to the trunk network.

DeepONet – Error Decomposition

Given **fixed parametrizations** $\mathcal{P}_1, \dots, \mathcal{P}_M$ and **evaluation points** $\hat{x}_1, \dots, \hat{x}_N$, the reference targets form the matrix

$$A = \begin{bmatrix} u_{\mathcal{P}_1}(\hat{x}_1) & \cdots & u_{\mathcal{P}_M}(\hat{x}_1) \\ \vdots & \ddots & \vdots \\ u_{\mathcal{P}_1}(\hat{x}_N) & \cdots & u_{\mathcal{P}_M}(\hat{x}_N) \end{bmatrix}$$

and the outputs of the DeepONet can be written as

$$\tilde{A} = \underbrace{\begin{bmatrix} t_1(\hat{x}_1) & \cdots & t_m(\hat{x}_1) \\ \vdots & \ddots & \vdots \\ t_1(\hat{x}_N) & \cdots & t_m(\hat{x}_N) \end{bmatrix}}_{=:T} \underbrace{\begin{bmatrix} b_1(\mathcal{P}_1) & \cdots & b_1(\mathcal{P}_M) \\ \vdots & \ddots & \vdots \\ b_m(\mathcal{P}_1) & \cdots & b_m(\mathcal{P}_M) \end{bmatrix}}_{=:B^\top} = [t_1 \quad \cdots \quad t_m] \begin{bmatrix} b_1^\top \\ \vdots \\ b_m^\top \end{bmatrix},$$

where B and T are discretizations of the branch and trunk net outputs. Similar to [Lanthaler et al. \(2022\)](#), we derive an **error decomposition into branch and trunk errors**

$$\mathcal{E} := \|A - \tilde{A}\|_F^2 = \underbrace{\|T(T^+ A - B^\top)\|_F^2}_{=: \mathcal{E}_B} + \underbrace{\|(I - TT^+)A\|_F^2}_{=: \mathcal{E}_T},$$

where $\|\cdot\|_F$ is the **Frobenius norm** and T^+ is the **pseudoinverse** of T .

Analyzing the DeepONet Using the Singular Value Decomposition

A **singular value decomposition (SVD)** of the target matrix A (rank k) reads as follows:

$$A = \Phi \Sigma V^\top = [\phi_1 \ \cdots \ \phi_k] \begin{bmatrix} \sigma_1 & 0 & \cdots & 0 \\ 0 & \sigma_2 & \cdots & 0 \\ \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \cdots & \sigma_k \end{bmatrix} \begin{bmatrix} v_1^\top \\ \vdots \\ v_k^\top \end{bmatrix} = \sum_{j=1}^k \sigma_j \phi_j v_j^\top,$$

where Φ and V are semi-orthogonal matrices. The **Eckart–Young theorem** states that the **best low-rank approximation of a matrix in the Frobenius norm** is given by truncating its SVD.

The structure is the same for the discretized DeepONet

$$\tilde{A} = [t_1 \ \cdots \ t_m] \begin{bmatrix} b_1^\top \\ \vdots \\ b_m^\top \end{bmatrix}.$$


Therefore, on the training data, we can minimize the discrete approximation error on the training data when choosing:

$$t_i = \phi_i \quad \text{and} \quad b_i = \sigma_i v_i.$$

DeepONet – Error Decomposition Results for the KdV Equation

Results for the **KdV equation** with $t_0 = 0.0$ and $\hat{t} = 0.2$.

900 training and 100 test configurations.



total error
 δ

trunk error
 δ_T

branch error
 δ_B

SVD truncation error
 δ_{SVD}

DeepONet – Branch Error

Using the left singular vectors, the **branch error** becomes

$$\mathcal{E}_B = \sum_{i=1}^m \sigma_i^2 \underbrace{\|b_i - v_i\|_2^2}_{=:L_i}.$$


We call

$$\sigma_i^2 L_i$$

the **(weighted) mode loss** because it equals the loss contribution of the i th mode. Accordingly, L_i is the **unweighted mode loss**.

This choice of **left singular vectors as the trunk basis** is often denoted **POD-DeepONet** in **Lu et al. (2022)**.

Unweighted mode loss



DeepONet – Branch Error

Using the left singular vectors, the **branch error** becomes

$$\mathcal{E}_B = \sum_{i=1}^m \sigma_i^2 \underbrace{\|b_i - v_i\|_2^2}_{=:L_i}.$$


We call

$$\sigma_i^2 L_i$$

the **(weighted) mode loss** because it equals the loss contribution of the i th mode. Accordingly, L_i is the **unweighted mode loss**.

This choice of **left singular vectors as the trunk basis** is often denoted **POD-DeepONet** in **Lu et al. (2022)**.

Unweighted mode loss



DeepONet – Branch Error

Using the left singular vectors, the **branch error** becomes

$$\mathcal{E}_B = \sum_{i=1}^m \sigma_i^2 \underbrace{\|b_i - v_i\|_2^2}_{=:L_i}.$$


We call

$$\sigma_i^2 L_i$$

the **(weighted) mode loss** because it equals the loss contribution of the i th mode. Accordingly, L_i is the **unweighted mode loss**.

This choice of **left singular vectors as the trunk basis** is often denoted **POD-DeepONet** in **Lu et al. (2022)**.

Unweighted mode loss



DeepONet – Branch Error

Using the left singular vectors, the **branch error** becomes

$$\mathcal{E}_B = \sum_{i=1}^m \sigma_i^2 \underbrace{\|b_i - v_i\|_2^2}_{=:L_i}.$$


We call

$$\sigma_i^2 L_i$$

the **(weighted) mode loss** because it equals the loss contribution of the i th mode. Accordingly, L_i is the **unweighted mode loss**.

This choice of **left singular vectors as the trunk basis** is often denoted **POD-DeepONet** in **Lu et al. (2022)**.

Unweighted mode loss



DeepONet – Branch Error

Using the left singular vectors, the **branch error** becomes

$$\mathcal{E}_B = \sum_{i=1}^m \sigma_i^2 \underbrace{\|b_i - v_i\|_2^2}_{=:L_i}.$$


We call

$$\sigma_i^2 L_i$$

the **(weighted) mode loss** because it equals the loss contribution of the i th mode. Accordingly, L_i is the **unweighted mode loss**.

This choice of **left singular vectors as the trunk basis** is often denoted **POD-DeepONet** in **Lu et al. (2022)**.

Unweighted mode loss



→ The coefficients of modes with **large singular values** are **learned best**. Errors for modes with **small singular values** **remain high**.

DeepONet – Branch Error

Using the left singular vectors, the **branch error** becomes

$$\mathcal{E}_B = \sum_{i=1}^m \sigma_i^2 \underbrace{\|b_i - v_i\|_2^2}_{=:L_i}.$$

We call


$$\sigma_i^2 L_i$$

the **(weighted) mode loss** because it equals the loss contribution of the i th mode. Accordingly, L_i is the **unweighted mode loss**.

This choice of **left singular vectors as the trunk basis** is often denoted **POD-DeepONet** in **Lu et al. (2022)**.

(Weighted) mode loss

4000 Training Steps



DeepONet – Branch Error

Using the left singular vectors, the **branch error** becomes

$$\mathcal{E}_B = \sum_{i=1}^m \sigma_i^2 \underbrace{\|b_i - v_i\|_2^2}_{=:L_i}.$$

We call


$$\sigma_i^2 L_i$$

the **(weighted) mode loss** because it equals the loss contribution of the i th mode. Accordingly, L_i is the **unweighted mode loss**.

This choice of **left singular vectors as the trunk basis** is often denoted **POD-DeepONet** in **Lu et al. (2022)**.

→ Analyzing the actual error contributions, the **modes with medium singular values contribute most**.

(Weighted) mode loss



DeepONet – Branch Error

Using the left singular vectors, the **branch error** becomes

$$\mathcal{E}_B = \sum_{i=1}^m \sigma_i^2 \underbrace{\|b_i - v_i\|_2^2}_{=:L_i}.$$


We call

$$\sigma_i^2 L_i$$

the **(weighted) mode loss** because it equals the loss contribution of the i th mode. Accordingly, L_i is the **unweighted mode loss**.


This choice of **left singular vectors as the trunk basis** is often denoted **POD-DeepONet** in **Lu et al. (2022)**.

(Weighted) mode loss



How to improve the performance on medium-sized singular value modes?

Finite Basis DeepONets (FBDONs)



Howard, Heinlein, Stinis (in prep.)

Variants:

Shared-trunk FBDONs (ST-FBDONs)

The trunk net learns spatio-temporal basis functions. In ST-FBDONs, we use the **same trunk network for all subdomains**.

Stacking FBDONs

Combination of the **stacking multifidelity approach** with FBDONs.

Heinlein, Howard, Beecroft, Stinis (2025)


FBDONs – Wave Equation

Wave equation

$$\frac{d^2s}{dt^2} = 2 \frac{d^2s}{dx^2}, \quad (x, t) \in [0, 1]^2$$

$$s_t(x, 0) = 0, x \in [0, 1], \quad s(0, t) = s(1, t) = 0,$$

Solution: $s(x, t) = \sum_{n=1}^5 b_n \sin(n\pi x) \cos(n\pi\sqrt{2}t)$



Parametrization

Initial conditions for s parametrized by $b = (b_1, \dots, b_5)$ (normally distributed):

$$s(x, 0) = \sum_{n=1}^5 b_n \sin(n\pi x) \quad x \in [0, 1]$$

Training on 1 000 random configurations.

Mean rel. ℓ_2 error on 100 config.

	Mean rel. ℓ_2 error on 100 config.
DeepONet	0.30 ± 0.11
ML-ST-FBDON ([1, 4, 8, 16] subd.)	0.05 ± 0.03
ML-FBDON ([1, 4, 8, 16] subd.)	0.08 ± 0.04

→ Sharing the trunk network does not only save in the number of parameters but even yields **better performance**

Cf. [Howard, Heinlein, Stinis \(in prep.\)](#)

Summary

Multilevel Finite Basis Physics Informed Neural Networks (ML-FBPINNs)

- Schwarz domain decomposition architectures **improve the scalability of PINNs** to large domains / high frequencies, **keeping the complexity of the local networks low.**
- As classical domain decomposition methods, **one-level FBPINNs** are **not scalable to many subdomains**; multilevel FBPINNs **enable scalability**.

Extensions to Stacking Multifidelity PINNs, RaNNs, and DeepONets

- Multifidelity stacking PINNs with FBPINNs improve **accuracy and efficiency** for time-dependent problems.
- RaNNs reduce computational cost but face **ill-conditioning**, mitigated by **Schwarz preconditioning** and **SVD**.
- DeepONets provide **efficient predictions** for parametrized problems but struggle with multiscale problems. Domain decomposition **improves scalability and performance**.

Thank you for your attention!



Topical Activity
Group
Scientific Machine
Learning

